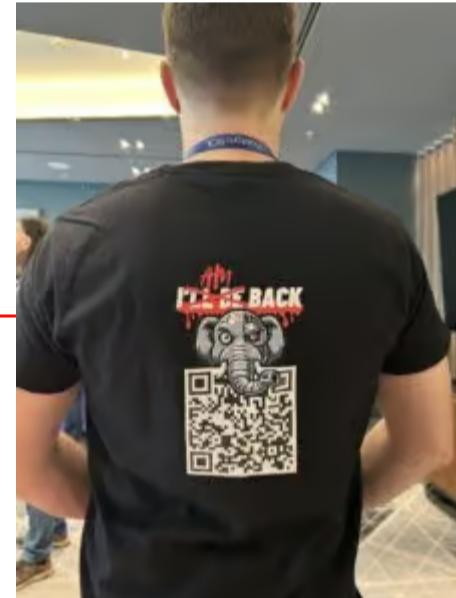


13 MAY 2025

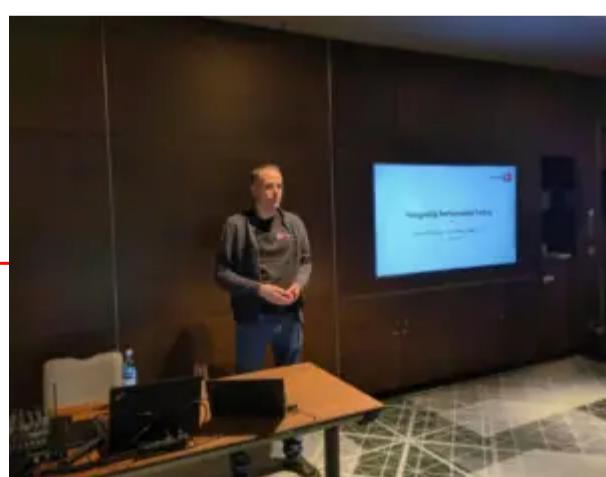
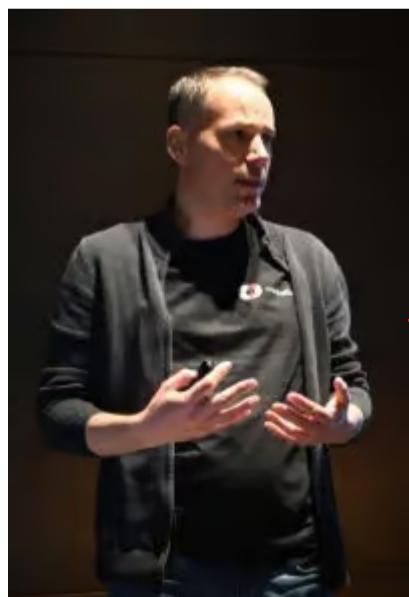
PostgreSQL Conference Germany 2025

[PGConf.DE 2025](#), the 9th Annual PostgreSQL Conference Germany, was held on May 8–9, 2025, at the Marriott Hotel near Potsdamer Platz in Berlin. The event interconnected PostgreSQL enthusiasts, developers, DBAs, and industry sponsors for two days of fascinating talks across four parallel tracks. It was the biggest event so far, with 347 attendees. The whole conference was very well organized, and therefore special thanks are due to all the organizers—in particular Andreas Scherbaum, the main organizer—for their efforts and hard work.

Our company, credativ GmbH, [being independent again](#), participated as a gold sponsor. The credativ CTO Alexander Wirt, Head of Sales & Marketing Peter Drew and team leader of Database team Tobias Kauder, were available for attendees at the credativ booth. Many thanks to our team colleague Sascha Spettmann for delivering all the stands and billboards to the conference and back again.



In total, we held four talks at the conference. Michael Banck, technical leader of our database team, presented the German-language talk "[PostgreSQL Performance Tuning](#)". He provided a deep and comprehensive overview of the most important performance-tuning parameters in PostgreSQL and explained how they influence the database's behavior. His talk attracted a large audience and was very well received.



I had an absolutely unique opportunity to present three different talks in the English track. In my regular talk “[PostgreSQL Connections Memory Usage: How Much, Why and When](#)”[↗], I presented the results of my research and tests on PostgreSQL connections’ memory usage. After explaining the most important aspects of Linux memory management and measurements of memory usage reported by standard commands, I detailed PostgreSQL connection memory usage during query execution based on numbers reported in smaps files. I intend to publish detailed blog posts about my findings soon. My other talk, “[Building a Data Lakehouse with PostgreSQL](#)”[↗], was originally chosen as a reserve talk in case of a last-minute cancellation. Unfortunately, this indeed happened: the talk “Creating a Board Game Chatbot with Postgres, AI, and RAG” by Matt Cornillon had to be replaced. The speaker could not attend because his flight was unexpectedly canceled at very short notice.



In the sponsor track, credativ CTO Alexander Wirt and I presented a talk titled “[Your Data Deserves the Best: Migration to PostgreSQL](#)”[↗]. It featured our new migration tool, “credativ-pg-migrator.” It is capable of migrating data models (tables, data, indexes, constraints, and views) from Informix, IBM DB2 LUW, MS SQL Server, Sybase ASE, SQL Anywhere, and MySQL/MariaDB. In the case of Informix, it can also convert stored procedures, functions, and triggers into PostgreSQL PL/pgSQL. We will share more details about this tool in a separate blog post.



Since there were always three or four parallel talks, I had to carefully choose which sessions to attend. I greatly enjoyed the talk “[András in Windowsland – a DBA’s \(mis\)adventures](#)”[↗] by András Váczi from Cybertec. The speaker presented many useful tips for accessing and troubleshooting PostgreSQL on Windows. I also enjoyed the German talk “[Modern VACUUM](#)”[↗], delivered by Christoph Berg from Cybertec. He provided valuable insights into the history and implementation details of the VACUUM command and autovacuum background processes. Another very interesting talk was the German presentation “[Modernes SSL ideal einsetzen](#)”[↗] by Peter Eisentraut from EDB. The talk covered selecting appropriate protocol versions and cryptographic cipher suites, managing keys and certificates, and configuring client/server settings to meet contemporary security standards. The talk “[Comparing the Oracle and PostgreSQL transaction systems](#)”[↗], delivered by Laurenz Albe from Cybertec, received a lot of well-deserved attention. Key topics included Oracle’s undo/redo mechanism versus PostgreSQL’s MVCC approach, differences in isolation level defaults and anomalies, and how each database implements constraints and transactional DDL. Last but not least, I want to mention the talk “[What is an SLRU anyway?](#)”[↗] delivered by major PostgreSQL contributor Álvaro Herrera. He explained that SLRUs are essentially circular logs with an in-memory cache used for tracking information such as committed transactions or snapshot data and he highlighted the significance of PostgreSQL 17’s innovations which made SLRU cache sizes configurable. Unfortunately, the talks were not recorded, but slides for the majority of the talks are already available on the conference website.

The whole event was highly informative and provided excellent networking opportunities. We are very much looking forward to participating in the next PGConf.DE. In the meantime, stay tuned to all credativ news, follow us on social media – [LinkedIn](#)  and [Mastodon](#) .

If you are interested in our PostgreSQL related services, click here!

JM

ABOUT THE AUTHOR

Josef Machytka

[View posts](#) 