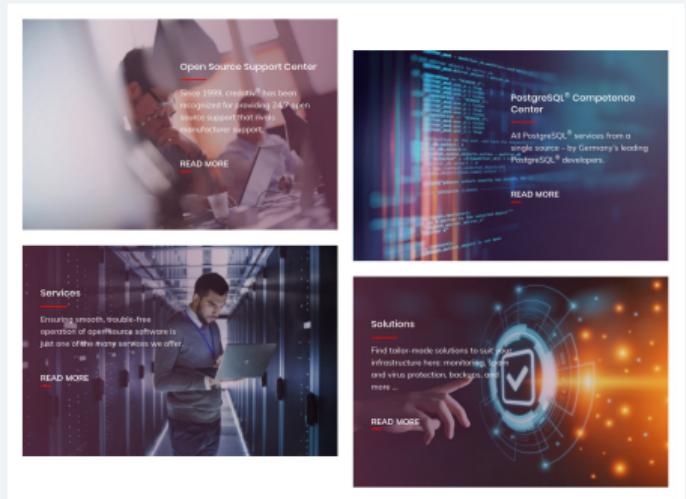


Gin, Btree_gin, GiST, Btree_gist, Hash and Btree indexes on JSONB data

Josef Machytka <josef.machytka@credativ.de>

2025-06-20 - credativ Tech Talk

- Founded 1999 in Jülich, Germany
- Close ties to Open-Source Community
- More than 40 Open-Source experts
- Consulting, development, training, support (3rd-level / 24x7)
- Open-Source infrastructure with Linux, Kubernetes, Proxmox
- Open-Source databases with PostgreSQL
- DevSecOps with Ansible, Puppet, Terraform and others
- Since 2025 independent owner-managed company again



- Professional Service Consultant - PostgreSQL specialist at credativ GmbH
- 30+ years of experience with different databases
- PostgreSQL (13y), BigQuery (7y), Oracle (15y), MySQL (12y), Elasticsearch (5y), MS SQL (5y)
- 10+ years of experience with Data Ingestion pipelines, Data Analysis, Data Lake and Data Warehouse
- 2+ years of practical experience with different LLMs / AI including their architecture and principles
- From Czechia, living now 11 years in Berlin

-  linkedin.com/in/josef-machytka
-  medium.com/@josef.machytka
-  youtube.com/@JosefMachytka
-  github.com/josmac69/conferences_slides
-  researchgate.net/profile/Josef-Machytka
-  academia.edu/JosefMachytka
-  sessionize.com/josefmachytka

All My Slides:



Recorded talks:



Table of contents

- Problems with implementation
- What was tested
- GIN indexes
- BTREE_GIN extension
- GIST & BTREE_GIST indexes
- HASH indexes
- BTREE indexes
- Decomposition of JSON data
- Statistics

Problems with implementation of JSON data

- Frontend and backend developers love the flexibility of JSON.
- JSON minimizes the need for app changes due to schema changes.
- IoT devices use JSON - W3C Web of Things Working Group standardized JSON for IoT.

- Data quality checks - absolute freedom might be a big challenge.
- Problems with data cleansing and transformation.
- Business intelligence, ML, and reporting need structured and standardized data.

- But the full decomposition of JSON can be a complex and painful task.
- Databases must handle JSON data, there is no escape.

Clients struggle with implementing JSONB

- Articles are often too shallow, repeating documentation.
- Very trivial examples - create a table, insert 3 rows, try explain, celebrate.
- Even ChatGPT-4o is not helpful with deeper and more complicated topics.

- Clients develop with small inadequate datasets.
- Tests are often too simple, just guessing production use cases.
- PostgreSQL dev instance has inadequate configuration.

- Confusion about TOAST tables, compression, and storage.
- Doubts about design - partitions vs 1 big table.
- Developers are obsessed with forcing indexes.

What was tested

What was tested

- Different types of indexes for different use cases
- Different compression and storage methods
- One big table vs partitioned tables
- Influence of PostgreSQL settings - memory, costs
- Performance under different loads - multiple simultaneous sessions
- Influence of parallelism
- Influence of data distribution and selectivity
- User defined statistics on JSONB columns
- Full decomposition vs one big JSONB column
- Deep dive into internals of indexes. Analysis of code

Dataset for tests



- GitHub Archive events - www.gharchive.org
- Separate .gz files for each hour - YYYY-MM-DD-HH24.json.gz
- One big JSONB column with all the data

```
CREATE TABLE github_events (
    id SERIAL PRIMARY KEY NOT NULL,
    jsonb_data JSONB);
```

GitHub events - JSON record



```
{  "id": "26167585827",
  "repo": {    "id": 581592468,
    "url": "https://api.github.com/repos/tiwabs/tiwbts_audio_door_tool",
    "name": "tiwabs/tiwbts_audio_door_tool" },
  "type": "PushEvent",
  "actor": {    "id": 48737497,
    "url": "https://api.github.com/users/tiwbts",
    "login": "tiwabs",
    "avatar_url": "https://avatars.githubusercontent.com/u/48737497?",
    "gravatar_id": "",
    "display_login": "tiwabs" },
  "public": true,
  "payload": {    "ref": "refs/heads/master",
    "head": "3ca247941f269bcdeb17e5b12e9b3b74b1c4da2",
    "size": 1,
    "before": "0dd5471667b12084b8fc88b1bca299780382d50a",
    "commits": [
      {
        "sha": "3ca247941f269bcdeb17e5b12e9b3b74b1c4da2",
        "url": "https://api.github.com/repos/tiwbts/....12e9b3b74b1c4da2",
        "author": { "name": "Tiwabs", "email": "mrskielz@gmail.com" },
        "message": "fix(export): export nametable if export succed",
        "distinct": true
      ],
      "push_id": 12149772587,
      "distinct_size": 1
    },
    "created_at": "2023-01-01T13:39:55Z" }
```


GitHub events - testing details

- Tested in PostgreSQL 15 and 16
- Python scripts for downloading, importing, analyzing, and testing
- Multiple local and AWS RDS testing environments
- Different CPUs, all with 8 cores and 32 GB RAM
- Used 1 week of data from January 2023
- In total 17,474,101 rows
- 3 tables, different compression methods:
 - pgz: 41 GB
 - lz4: 38 GB
 - external storage with no compression: 98 GB



Performance testing - sequential scan on the table

- Aggregation query over all records using sequential scan on the table, without parallelism
- The old compression method, pglz, was already slower than no compression with 8 sessions on 8 cores
- With 16, 32, and 64 sessions on 8 cores, pglz became a serious performance bottleneck



Disk intensive operations – throughput matters

- The same problem occurs on ALL clouds; we just tested it on AWS
- On AWS RDS SSD 300GB with 3,000 IOPS, the throughput of 125 MiBps was a real disaster
- All disk-intensive operations were 4x to 5x slower than on the local PC
- With SSD 500GB and 12,000 IOPS, and a throughput of 500 MiBps, we finally got reasonable results
- But auto-scaling of the disk can also slow down your actions by 5x or more
- Never try to save money on a cloud instance by using a slow, small disk
- Do not rely only on the auto-scaling of the disk

GIN indexes

- Showed very stable performance even under high load
- But for their usage proper settings are crucial
- Shared_buffers -> 25% of RAM, effective_cache_size -> 50% of RAM
- GIN indexes do not support parallelism, neither for creation nor for usage
- Parallelism can be a significant factor in using or not using GIN indexes
- If parallel workers are available, the planner can choose parallel sequential scan on the table
- If all parallel workers are in use, the planner uses GIN indexes for new queries
- Set Max_parallel_workers_per_gather = 0 at least for the query


- SSD: random_page_cost = 1.1, effective_ioConcurrency = 200
- Set random_page_costs <= seq_page_cost (=1) if the database is fully cached in memory
- Different values of work_mem had minimal impact if the query used GIN index scan
- PostgreSQL code: `src/backend/optimizer/path/costsize.c`
 - `cpu_tuple_cost`, `cpu_index_tuple_cost`, `cpu_operator_cost` - how they are measured?
 - `parallel_setup_cost`, `parallel_tuple_cost` - how are these measured?
 - The code says "measured on an arbitrary scale"
 - Especially `cpu_tuple_cost` is used incredibly often in the code
 - Its value influences the planner's decisions significantly

- It can take hours to create a new GIN index on the whole column with existing data
- `maintenance_work_mem` had only small impact on the speed of creating a GIN index
- Disk IO is the main factor affecting the speed of creating a GIN index


- Updates of GIN indexes become significantly slower as the table size grows
- The index is rebuilt when the `gin_pending_list_limit` is reached or during vacuuming
- Default value of `gin_pending_list_limit` is 4MB = 512 data pages

- The size of the table matters
- The speed of inserting rows per second can decrease by up to 50%
- Partitioning can help significantly. However, disk IO is again the main factor


GIN indexes – speed of inserts – one big table




GIN indexes – speed of inserts – partitions



GIN indexes – creation – total progress



GIN indexes – creation – difference in progress



Gin Indexes inspection - pending pages and tuples



- We can use extensions to get some deeper information about GIN indexes
- pgstattuple:
 - pgstatginindex()
- pageinspect:
 - gin_page_opaque_info() - basic info about page
 - gin_metapage_info() - details for metapage
 - gin_leafpage_items() - details for leaf page

```
SELECT * FROM pgstatginindex('index_name');
```

version	pending_pages	pending_tuples
2	414	1853

```
SELECT *
FROM gin_metapage_info(
    get_raw_page('index_name', 0))\gx
```

```
pending_head      : 292675
pending_tail      : 339992
tail_free_size   : 220
n_pending_pages  : 414
n_pending_tuples : 1853
n_total_pages    : 339200
n_entry_pages    : 312283
n_data_pages     : 24533
n_entries        : 52572205
version          : 2
```


--> but before VACUUM these
values are only estimates!

Gin Indexes inspection - deeper dive into pages stats



```
-- How to get proper count of pages?  
pg_class: 339986, metapage: 339200 - both are estimates, just taken differently  
  
-- Let's calculate the proper count of pages from the size of data files  
  
SELECT pg_relation_size('index_name') / 8192;  
-> 357105 pages  
  
-- Now we can get statistics about GIN index pages  
  
WITH pages AS (  
    SELECT *  
    FROM generate_series(0,  
        (SELECT pg_relation_size('index_name') / 8192) -1) as pagenum)  
SELECT  
    (SELECT flags  
    FROM gin_page_opaque_info(  
        get_raw_page('index_name', pagenum))) as flags,  
    count(*) as pages  
FROM pages GROUP BY flags ORDER BY flags;
```

GIN indexes – rebuild of index during insertion of data



- GIN index with jsonb_ops operator class is the most versatile but also the biggest
- It allows searching for equality of values on multiple unknown levels of keys
- The @? and @@ operators can be used with * and ** wildcards
- Example: `WHERE jsonb_data @@ '$.*' == "python3"`
- The size of the jsonb_ops GIN index on the whole column can reach 80% of the table size

- The operator class jsonb_path_ops works only with fully known jsonpath
- It allows searching for equality of values on multiple known levels of keys
- The @? and @@ operators cannot use wildcards, the jsonpath must be known
- Example: `WHERE jsonb_data @@ '$.payload.pull_request.head.repo.topics[*]' == "python3"`
- The GIN index with jsonb_path_ops on the whole column can reach 30% of the table size


- If the second object is contained in the first one - an exact match of the key(s) and value(s)
- Works with both operator classes
- Works for nested objects and arrays
- Allows searching for equality of multiple values in one condition
- Searching for values from lists of values - events from specific users, a specific repository
- Run times are in dozens or hundreds of milliseconds
- Very stable performance even with multiple sessions running in parallel
- Limitation - the path must be known
- This will find data: `WHERE jsonb_data @> '{"payload":{"commits":[{"author":{"name": "Jane Joy"}]}]}`
- This will not find: `WHERE jsonb_data @> '{"commits":[{"author":{"name": "Jane Joy"}]}]`

- Operators `? , ?| , and ?&`
- They are used to look for the existence of key(s) on the top level
- These operators only work with the `jsonb_ops` operator class
- The usage of the GIN index depends on statistics
- If a key is present in the majority of records, the GIN index is not used
- If the table is very small, the GIN index is not used
- The GIN index is only used for keys that are not present in the majority of records
- Useful for a very dynamic schema or a table that stores many different JSON datasets

- SQL\JSON contains multiple amazing methods, but GIN index does not work for them.
- `like_regex` - tests if the string value returned by jsonpath matches a regular expression:
`WHERE jsonb_data @? '$.description ? (@ like_regex ".*Michigan.*")'`
- `starts with` - tests if the string value returned by jsonpath starts with a specific string:
`WHERE jsonb_data @? '$.laureates[*].firstname ? (@ starts with "Jo")'`
- `exists` - tests if a key exists in the JSONB schema at a given level:
`WHERE jsonb_data @? '$.laureates[*].firstname ? (exists (@))'`
- The PostgreSQL community should consider creating indexes for these operators

- GIN index with tsvector_ops operator class allows full text search
- The function jsonb_to_tsvector converts JSONB data into tsvector
- Example: `WHERE jsonb_to_tsvector('english', jsonb_data, "string") @@ to_tsquery('search_string')`
- Full text search works for equality of words/synonyms
- You can combine words using AND/OR
- The tsvector_ops index on the whole column can be larger than the table
- It only makes sense to create an index on free text columns
- It speeds up search by at least 100 times
- Performance is very stable under high load

GIN indexes – full text search – commit messages



- Gin_trgm_ops operator class allows string search using LIKE
- The index over the whole column does not distinguish keys and values
- It still performs an equality search behind the scenes - equality of trigrams
- Creating an index on free text columns is the only scenario where it makes sense

- The size of the gin_trgm_ops GIN index on the whole column can reach 50% of the table size
- It significantly speeds up search, even up to 1000x
- The performance is very stable under high load

- Partitioned tables showed multiple advantages over one big table
- Query run times using GIN indexes are faster on partitioned tables, approximately 5 times faster
- Loading data into partitioned tables is faster
- Updates of GIN indexes on partitions are faster
- Partitioning also makes GIN indexes more efficient in case of very uneven data distribution
- Bloat of partial GIN indexes due to the values with very high frequency is much lower

BTREE_GIN indexes

- The BTREE_GIN extension combines the BTREE and GIN indexes
 - It adds GIN operator classes with BTREE behavior
 - Any GIN operator class works with the BTREE_GIN index
 - The BTREE_GIN index can have multiple columns
 - It will optimize the search for any combination of these columns
 - The order of columns does not seem to be important
- The runtime with the BTREE_GIN index was better than with the GIN index + filter search
- The run times of the tested use cases were in the range of hundreds of milliseconds
- The performance was stable even with many parallel sessions

GIST indexes

- For indexing geo data, you need GIST indexes
- Most commonly in GeoJSON format
- Usually - Type (Point), coordinates [longitude (+/- 0-180), latitude (+/- 0-90)]

```
-- NASA meteorites dataset
{ "id": "1",
  "fall": "Fell",
  "mass": "21",
  "name": "Aachen",
  "year": "1880-01-01T00:00:00.000",
  "reclat": "50.775000",
  "reclong": "6.083330",
  "nametype": "Valid",
  "recclass": "L5",
  "geolocation": {
    "type": "Point",
    "coordinates": [ 6.08333, 50.775 ] } }
```

GIST indexes – PostGIS example

- Let's create a GIST index based on GEOMETRY(point, 4326) PostGIS data type.
- EPSG code 4326 is for WGS 84 spacial reference system.

```
-- we can create a GIST index on a GEOMETRY column manually:
```

```
CREATE INDEX ON nasa_meteorits USING GIST(
    ST_SETSRID(ST_MakePoint(
        cast(jsonb_data->'geolocation'->'coordinates'->>0 as float),
        cast(jsonb_data->'geolocation'->'coordinates'->>1 as float) ), 4326) );
```

```
-- or use PostGIS extension function st_geomfromgeojson
```

```
-- expects a GeoJSON object as input, recognizes content automatically:
```

```
-- meteorites: { "type": "Point", "coordinates": [ 6.08333, 50.775 ] }
```

```
-- earthquakes: { "geometry": { "type": "Point", "coordinates": [ -104.024, 31.646, 6.8514 ] } }
```

```
CREATE INDEX ON nasa_meteorits USING GIST(
    ST_GeomFromGeoJSON(jsonb_data->'geolocation') );
```

- The BTREE_GIST extension allows to combine GIST and BTREE indexes
 - We cannot create a GIST index on a whole JSONB column
 - We can combine multiple columns into a BTREE_GIST with different operator classes
-
- intarray extension: `gist_int_ops` and `gist_intbig_ops` operator classes for arrays
 - `gist_trgm_ops` operator class for performing LIKE search over strings
 - `tsvector_ops` operator class for creating a GIST index for full-text search

- Earthquakes dataset - United States Geological Survey (earthquake.usgs.gov)
- GIST index on JSONB column combining multiple extracted values
- Geolocation, magnitude as a number, place as a trigram, and magnitude type as a list of values
- Optimizes all variants of queries using these columns
- Quick to create - 1 minute on a 1 GB dataset. Size is 20% of the table size

```
CREATE INDEX ON jsonimport USING gist (
    ST_GeomFromGeoJSON(jsonb_data->'geolocation'),
    ((jsonb_data->'properties'->>'mag')::numeric),
    (jsonb_data->'properties'->>'place') gist_trgm_ops,
    (jsonb_data->'properties'->>'magType') );
```


HASH indexes

- HASH indexes allow only equality search, not range search
- Support only single-column indexes and do not allow uniqueness checking
- Can be smaller and quicker than BTREE indexes, but only on some data
- In my tests, the performance of HASH and BTREE indexes on unique data was almost the same
- On non-unique data, HASH indexes were still very quick, but at least 5x slower than BTREE indexes
- However, HASH indexes have been bigger in size than BTREE indexes - 2x to 3x
- HASH indexes are extremely sensitive to the data distribution
- Require unique or nearly unique data with a low number of rows per hash bucket
- Many repeated values lead to a big number of overflow bucket pages and bad performance

- Always check PostgreSQL statistics on the data before deciding to use a HASH index
- For JSONB data, create user defined statistics for the keys you want to index
- In pg_stats, a value of n_distinct = -1 indicates fully unique values
- If n_distinct != -1, check the values in most_common_vals and most_common_freqs
- If the values in most_common_freqs are 0.1 or higher, a HASH index is not a good choice
- You can avoid NULLs and empty strings by using a partial index

HASH index - creation / reindex

- Numbers are processed quicker than strings
- Unique data is processed very quickly
- Non-unique strings with many duplicates are processed extremely slowly
- Unique integers: 30 seconds,
highly duplicate integers: 1.5 minute
- Unique strings: 1 minute,
highly duplicate strings: 55 minutes



- Extension `pageinspect` contains functions for inspecting the index
- Function `hash_page_type` returns the type of the page
- If all values are unique, index contains: metapage (0), bucket pages, bitmap page (last)
- If there are duplicates, overflow pages are added after the bitmap page
- The bitmap pages maintain the map of free and in-use overflow pages
- Free overflow pages, not attached to any bucket, are kept and reused
- There are other functions to inspect content of pages
- `Hash_metapage_info`, `hash_bitmap_info`, `hash_page_items`, `hash_page_stats`

HASH indexes - maintenance

- HASH indexes are not automatically maintained
- Empty overflow pages are never removed from the index
- Shrinking requires REINDEX or VACUUM FULL

- HASH indexes can be useful only in very specific cases
- BTREE indexes seem to be a much better choice

BTREE indexes

- Very small and quick -> ideal first choice
- Allow parallel index build and scan
- Can be created in minutes, even on large tables
- Support equality and range queries: <, <=, =, >=, and >
- With `text_pattern_ops` (for each column), can be used for prefix-LIKE queries
- Some transformations must be encapsulated into immutable functions
- Conditions in queries must contain the exact indexed expression
- Partial BTREE indexes can be very useful for dynamic schemas
- LIMIT improves delivery of results significantly

Sizes of indexes

Summary of results – sizes of indexes

Table - lz4 TOAST compression, 17.5 M rows	38 GB	
GIN index - jsonb_ops - whole JSONB column	25 GB	66 %
GIN index - jsonb_path_ops - whole JSONB column	16 GB	42 %
GIN index - gin_trgm_ops - whole JSONB column	16 GB	42 %
GIN index - tsvector_ops - jsonb_to_tsvector, "string" values	34 GB	90 %
GIN index - tsvector_ops - just commit messages	0.5 GB	1.5 %
GIN index - gin_trgm_ops - just commit messages	1 GB	3 %
BTREE_GIN index - 'payload' jsonb_ops + created_at	23 GB	60 %
BTREE_GIN index - 'payload' jsonb_path_ops + created_at	15 GB	40 %
BTREE index on "created_at" timestamp	120 MB	0.2 %
HASH index on "created_at" timestamp	0.5 GB	1.5 %

JSON decomposition

- GitHub events schema has 936 unique jsonpaths
- Some keys contain different data types in different records
- Schema has 12 embedded arrays with JSONB objects -> additional 12 tables
- Main table would need 807 columns (jsonpaths without array elements)
- In total: text 604, number 105, boolean 70, datetime 28 columns
- I tried the main table: "ERROR: row is too big: size 9088, maximum size 8160"
- Some big nested JSONB objects would need separate tables too
- Decomposition is huge task on its own, prone to many errors

- A table with hundreds of columns is hard to use
- The theoretical limit is 1600 columns in the tuple
- But, the tuple must fit into one data page (8KB)

- Full jsonpath as a column name can easily exceed 63 characters
- Table with many columns requires careful design due to data types padding
- Columns must fit into 8-byte blocks - a 64-bit CPU reads a block of 8 bytes
- Alignment/padding can waste space for many small columns

- Nested composed data types can make the solution even more complex
- They use extended storage, i.e. TOAST
- This way you just convert one binary object into another
- Queries require encapsulation of top-level keys into parentheses
- Only after really trying it you will realize how big challenge it can be

- Manual decomposition is not worth the trouble
- You can use GENERATED columns for some frequently used jsonpaths
- This way you avoid a lot of manual work

```
ALTER TABLE github_events
ADD COLUMN actor_login text
GENERATED ALWAYS AS ((jsonb_data->'actor'->>'login')) STORED;
```

Statistics

Statistics for the whole JSONB object



- The planner seems to be able to deduce statistics for top-level keys.
 - For specific jsonpaths use CREATE STATISTICS command.
 - This way you will have always up-to-date statistics for your use cases.
-
- Command CREATE STATISTICS only prepares statistics object.
 - Statistics are gathered by first ANALYZE command and later by system as usual.

Statistics for the whole JSONB object



```
CREATE STATISTICS github_actor_login
ON ( ((jsonb_data -> 'actor'::text) ->> 'login'::text) )
FROM github_events;

CREATE STATISTICS github_created_at_ts
ON ( json_datetime_to_timestamp((jsonb_data ->> 'created_at'::text)) )
FROM github_events;
```

- View pg_stats_ext_exprs shows statistics.
- Columns n_distinct, most_common_vals, and most_common_freqs, correlation as in pg_stats.

- Indexes are not the "silver bullet" for everything.
- Don't be obsessed with forcing PostgreSQL to use indexes.
- The usage of indexes depends on frequency, selectivity, and correlation.
- In some use cases, a parallel sequential scan can be better than an index scan.
- The runtime of queries depends on data distribution - sorting in memory vs on disk.
- On multi-tenant systems, things are even more complicated.
- Understand your data and use cases to use the right tools!
- Make sure that your PostgreSQL has proper settings.

Thank you for your attention!



All my slides



Recorded talks

