


PostgreSQL Native Partitioning

From Ancient Monolithic Tables
to Scalable Better Future

Josef Machytka <josef.machytka@credativ.de>

2025-08-27 - datativ Tech Talk

- Founded 1999 in Jülich, Germany
- Close ties to Open-Source Community
- More than 40 Open-Source experts
- Consulting, development, training, support (3rd-level / 24x7)
- Open-Source infrastructure with Linux, Kubernetes, Proxmox
- Open-Source databases with PostgreSQL
- DevSecOps with Ansible, Puppet, Terraform and others
- Since 2025 independent owner-managed company again



- Professional Service Consultant - PostgreSQL specialist at credativ GmbH
- 33+ years of experience with different databases
- PostgreSQL (13y), BigQuery (7y), Oracle (15y), MySQL (12y), Elasticsearch (5y), MS SQL (5y)
- 10+ years of experience with Data Ingestion pipelines, Data Analysis, Data Lake and Data Warehouse
- 3+ years of practical experience with different LLMs / AI / ML including architecture and principles
- From Czechia, living now 12 years in Berlin

-  linkedin.com/in/josef-machytka
-  medium.com/@josef.machytka
-  youtube.com/@JosefMachytka

-  github.com/josmac69/conferences_slides
-  researchgate.net/profile/Josef-Machytka
-  academia.edu/JosefMachytka
-  sessionize.com/josefmachytka

All My Slides:




Recorded talks:



Why we need Partitioning?

Partitioning in General

- Splits large tables into smaller, more manageable partitions
- Assigns rows to partitions using a defined partition key
- Better query and maintenance performance via smaller tables
- Scans and processes only relevant partitions for each query
- Reduces index size and speeds up index lookups
- Enables efficient parallel query execution across partitions
- Supports targeted backup and restore of individual partitions
- Allows fast data lifecycle management by dropping partitions



Picture from Oracle documentation

- PostgreSQL MVCC turns UPDATEs into DELETE + INSERT
- Large tables with frequent updates can bloat quickly
- Heap pages fill with dead tuples, slowing scans
- Changes are randomly distributed across the table
- Large index scans on big table cause massive random IO
- Indexes only know tuple CTIDs, not visibility
- Every scan must check visibility of all tuples in table
- Huge tables lead to slower VACUUM and ANALYZE operations
- VACUUM must scan entire table to reclaim space



Image from the article [Snapshot Isolation in PostgreSQL](#)

Clustering

- Physically orders table data on disk by specified column(s)
- Rows with similar values are stored together
- Improves range queries by reducing random disk IO
- Most effective for immutable datasets
- Widely used in Data Warehouses and Data Lakehouses
- Partitioning & Clustering are complementary techniques
- Partitioning divides data into smaller tables
- Clustering organizes data within those tables
- Together they enhance query performance and manageability

Order Date	Country	Status
2023-03-02	US	Shipped
2023-03-04	AU	Processing
2023-03-05	JP	Canceled
2023-03-06	IN	Processing
2023-03-02	IN	Shipped
2023-03-05	US	Canceled
2023-03-04	JP	Shipped
2023-03-04	IN	Shipped
2023-03-06	JP	Processing
2023-03-02	AU	Canceled
2023-03-05	JP	Canceled
2023-03-06	AU	Processing
2023-03-05	US	Shipped
2023-03-06	IN	Processing
2023-03-02	AU	Shipped
2023-03-04	US	Shipped

Order Date	Country	Status
2023-03-02	AU	Canceled
2023-03-02	AU	Processing
2023-03-06	AU	Processing
2023-03-04	AU	Shipped
2023-03-02	IN	Processing
2023-03-04	IN	Processing
2023-03-06	IN	Shipped
2023-03-06	IN	Shipped
2023-03-05	JP	Canceled
2023-03-04	JP	Canceled
2023-03-05	JP	Processing
2023-03-06	JP	Shipped
2023-03-05	US	Canceled
2023-03-02	US	Shipped
2023-03-05	US	Shipped
2023-03-04	US	Shipped



Non-clustered Table

Table Clustered by Country and Status

Image from the article [Google BigQuery Data Structure](#)

Typical Use Cases

- Time-series data - monthly, weekly, daily, hourly partitions
- Multi-tenant applications - separate partitions for each tenant
- Quick access to specific data - range or list partitioning
- Data collected from multiple sources - hash partitioning
- Data Warehouse & Data Lakehouse - multiple dimensions
- Data archiving & Data Retention Policies (GDPR) - time-based partitioning



Images from the article [Scaling IoT Time-Series Data in Postgres](#)

Worst Case Scenarios

- Production logs from thousands of microservices in single table
- Daily ingestion exceeded 1 TB with unpredictable spikes
- Multiple indexes required for analytics and reporting queries
- Retention policy kept recent data for several days
- Deletion of old data used time and complex filtering criteria
- High write and delete rates caused huge table and index bloat



Image from the article [Logging in Microservices](#)

Real Life Worst Case Example

- PostgreSQL splits large tables into 1 GB segments
- 1 TB table creates 1000+ segment files
- Multi-GB indexes span hundreds of additional files
- Traffic spikes cause unpredictable data volume surges
- Peak days see extreme spikes in insert rates
- Large deletes may run for hours
- Autovacuum must reclaim space in table and indexes
- Bloat severely degrades query and write speed




Image from the article [Vacuum in PostgreSQL](#)

- Performance issues are often blamed on PostgreSQL itself
- Technical debt accumulates when proper design is postponed
- Production systems rarely allow downtime for major changes
- Feature development takes priority over infrastructure fixes
- Teams hesitate without clear partitioning strategy guidance
- "Why doesn't PostgreSQL handle large tables automatically?"
- NoSQL databases are often seen as magical scaling solutions
- Reality: No database can compensate for poor design choices

Automatic Partitioning

Magic Comes with a Price

- Some commercial databases offer automatic partitioning features
 - Snowflake organizes tables into internal micro-partitions
 - Micro-partitioning is fully managed and not user-configurable
 - Oracle Autonomous Database provides automatic partitioning
 - Skips tables with manual partitioning and analyses others
 - Automatic analysis may consume significant system resources
 - Both solutions are proprietary and involve high licensing costs




Image from the article [Snowflake documentation: Micro-partitions & Data Clustering](#)

- Open source databases primarily support manual partitioning
- Semi-automatic tools exist, built on manual partitioning foundations
- Users define initial setup and partitioning strategy
- Tools automate partition creation and routine maintenance tasks
- PostgreSQL: pg_partman for general partition management
- Time-series data: TimescaleDB extension provides automation

How to Implement Partitioning Later?

- Retrofitting native partitioning requires careful planning
 - Existing table must attach as a partition to new parent
 - May require data migration or complete table rewrites
 - Table inheritance offers simpler retrofitting approach
 - Inheritance useful when partitioning strategy unclear
 - Create new partitions for future data with routing rules
 - Practical examples will be demonstrated in other talk
-
- [Crunchy Bridge Guide: Partitioning Existing Tables in PostgreSQL](#)

```
BEGIN
    ALTER TABLE probes_measurements
        RENAME TO probes_measurements_old_data;

    CREATE TABLE probes_measurements (
        id SERIAL PRIMARY KEY,
        -- all other columns ...
        inserted_at TIMESTAMP WITH TIME ZONE
            DEFAULT clock_timestamp()
    ) PARTITION BY RANGE
        (measurement_timestamp);

    ALTER TABLE probes_measurements
        ATTACH PARTITION probes_measurements_old_data
        FOR VALUES FROM ('2024-11-13') TO ('2024-11-14');

    CREATE TABLE probes_measurements_20241113
        PARTITION OF probes_measurements
        FOR VALUES FROM ('2024-11-13') TO ('2024-11-14');

END;
```

PostgreSQL Partitioning Overview

- PostgreSQL supports table inheritance and declarative partitioning
- Table inheritance is legacy, not recommended for new applications
- It is more flexible, but requires manual management
- Declarative partitioning automates maintenance operations
- Each method has specific use cases and operational trade-offs
- Partitioning behavior is influenced by configuration settings
- Performance degrades significantly with excessive partitions





Image from the article [Partitioned and Inherited Tables in PostgreSQL](#)


Partitioning: Scalability and Planning Overheads

- Performance can degrade with hundreds/thousands of partitions
- Query planning time increases significantly
- Planning can exceed execution time for complex queries
- Many partitions may cause session memory exhaustion
- Each partition adds metadata overhead in the planner's cache

Cold cache



Warm cache




Images from the article [Postgres experiment: number of partitions vs. planning time](#)

- Maintenance operations are more efficient on smaller partition indexes
 - GIN indexes on large, unpartitioned tables badly degrade performance
 - Indexes on partitions are smaller, enabling faster index scans
 - Partitioning can yield 5x–10x faster queries in practical benchmarks
-
- [Benchmarking PostgreSQL: The Hidden Cost of Over-Indexing](#)
 - [How Partial and Covering Indexes Affect Update Performance in PostgreSQL](#)

PostgreSQL Native Partitioning

- Parent table is defined with the `PARTITION BY` clause
- Partition key columns must be included in the `PRIMARY KEY`
- Parent table is virtual - contains no data, only metadata
- Specifies columns and constraints for all partitions
- Child partitions inherit all columns and constraints from parent
- Only leaf partitions on the lowest level physically store data
- Each partitioning level can define a `DEFAULT` partition
- DML operations are automatically routed to the correct partition



Queries entire table Queries only 2020 partition

Image from the article
[PostgreSQL Insider - Partitioning Overview](#)

- **Range partitioning:**


- Partitions data into non-overlapping intervals
- Ideal for time series, numeric ranges, or sequential data
- Each partition covers a range: lower bound inclusive, upper bound exclusive
- Example: 0 to 10 means: $0 \leq x < 10$

- **List partitioning:**

- Assigns explicit values to each partition
- Suitable for categorical data, such as ENUMs or specific IDs

- **Hash partitioning:**

- Distributes rows evenly using a hash function
- Each partition defined by modulus and remainder of the hash value
- Useful for balancing data when no natural range or list exists



Images from the article
[PostgreSQL Insider - Partitioning Overview](#)

Example of Range Partitioning

```
CREATE TABLE orders_range (
    order_id SERIAL,
    customer_id INT NOT NULL,
    order_date DATE NOT NULL,
    status TEXT NOT NULL
) PARTITION BY RANGE (order_date);

CREATE TABLE orders_range_2022 PARTITION OF orders_range
FOR VALUES FROM ('2022-01-01') TO ('2023-01-01');

CREATE TABLE orders_range_2023 PARTITION OF orders_range
FOR VALUES FROM ('2023-01-01') TO ('2024-01-01');

CREATE TABLE orders_range_2024 PARTITION OF orders_range
FOR VALUES FROM ('2024-01-01') TO ('2025-01-01');

CREATE TABLE orders_range_default PARTITION OF orders_range DEFAULT;
```

Example of List Partitioning

```
CREATE TABLE orders_list (
    order_id SERIAL,
    customer_id INT NOT NULL,
    order_date DATE NOT NULL,
    status TEXT NOT NULL
) PARTITION BY LIST (customer_id);

CREATE TABLE orders_list_customer_1 PARTITION OF orders_list
FOR VALUES IN (1);

CREATE TABLE orders_list_customer_2 PARTITION OF orders_list
FOR VALUES IN (2);

CREATE TABLE orders_list_customer_3 PARTITION OF orders_list
FOR VALUES IN (3);

CREATE TABLE orders_list_default PARTITION OF orders_list DEFAULT;
```

Example of Hash Partitioning

```
CREATE TABLE orders_hash (
    order_id SERIAL,
    customer_id INT NOT NULL,
    order_date DATE NOT NULL,
    status TEXT NOT NULL
) PARTITION BY HASH (customer_id);

CREATE TABLE orders_hash_customer_1 PARTITION OF orders_hash
FOR VALUES WITH (MODULUS 4, REMAINDER 0);

CREATE TABLE orders_hash_customer_2 PARTITION OF orders_hash
FOR VALUES WITH (MODULUS 4, REMAINDER 1);

CREATE TABLE orders_hash_customer_3 PARTITION OF orders_hash
FOR VALUES WITH (MODULUS 4, REMAINDER 2);

CREATE TABLE orders_hash_customer_4 PARTITION OF orders_hash
FOR VALUES WITH (MODULUS 4, REMAINDER 3);
```

- Partitions can be attached or detached without data movement
- Each partition can reside on a different tablespace for storage optimization
- Query planner prunes partitions using partition key predicates
- Partition-wise joins and aggregates enable parallel processing
- SPLIT PARTITION and MERGE PARTITION are not yet supported natively

```
ALTER TABLE orders_range
    ATTACH PARTITION orders_range_2021
    FOR VALUES FROM ('2021-01-01') TO ('2022-01-01');
```

```
ALTER TABLE orders_range
    DETACH PARTITION orders_range_2019;
```

- Columns & Constraints
 - Columns or constraints added to parent are propagated to all partitions
 - Dropping a column or constraint from parent removes it from partitions
- Indexes
 - Indexes created on parent are automatically created on all partitions
 - New partitions inherit all parent indexes at creation time
 - Attaching a table as partition creates any missing required indexes
 - Dropping an index on parent removes it from all partitions
 - Indexes created via parent cannot be dropped individually on partitions
 - Indexes created directly on a partition exist only on that partition

```
SELECT
    n.nspname                                AS schema_name,
    c.relname                                 AS table_name,
    pt.partstrat                             AS strategy,      -- 'r' = range, 'l' = list, 'h' = ↪
    hash
    pt.partnatts                            AS key_cols,
    pg_get_partkeydef(c.oid)                 AS partition_key
FROM pg_partitioned_table AS pt
JOIN pg_class      AS c ON c.oid = pt.partrelid
JOIN pg_namespace AS n ON n.oid = c.relnamespace
WHERE n.nspname NOT IN ('pg_catalog', 'information_schema')
ORDER BY 1, 2;
```

schema_name	table_name	strategy	key_cols	partition_key
public	orders_hash	h	1	HASH (customer_id)
public	orders_list	l	1	LIST (customer_id)
public	orders_range	r	1	RANGE (order_date)

Query - All children and their parent tables

```
SELECT
    pn.nspname                      AS parent_schema,
    p.relname                        AS parent_table,
    cn.nspname                       AS partition_schema,
    c.relname                         AS partition_table
FROM pg_inherits AS i
JOIN pg_class      AS p ON p.oid = i.inhparent
JOIN pg_namespace AS pn ON pn.oid = p.relnamespace
JOIN pg_class      AS c ON c.oid = i.inhrelid
JOIN pg_namespace AS cn ON cn.oid = c.relnamespace
WHERE p.relkind = 'p'              -- parent is a partitioned table
ORDER BY parent_schema, parent_table, partition_schema, partition_table;
```

parent_schema	parent_table	partition_schema	partition_table
public	orders_hash	public	orders_hash_customer_1
public	orders_hash	public	orders_hash_customer_2
public	orders_hash	public	orders_hash_customer_3
public	orders_hash	public	orders_hash_customer_4
public	orders_list	public	orders_list_customer_1
public	orders_list	public	orders_list_customer_2
public	orders_list	public	orders_list_customer_3
public	orders_list	public	orders_list_default
public	orders_range	public	orders_range_2022
public	orders_range	public	orders_range_2023
public	orders_range	public	orders_range_2024
public	orders_range	public	orders_range_default

Query - Inspect the partition tree

```
SELECT
    t.relid::regclass      AS rel_name,
    t.parentrelid::regclass AS parent_name,
    t.level,
    t.isleaf
FROM pg_partition_tree('public.orders_list'::regclass) AS t
ORDER BY t.level, rel_name;
```

rel_name	parent_name	level	isleaf
orders_list		0	f
orders_list_customer_1	orders_list	1	t
orders_list_customer_2	orders_list	1	t
orders_list_customer_3	orders_list	1	t
orders_list_default	orders_list	1	t

Query - Bounds for all partitions of all parent tables

```
SELECT
    child.oid::regclass                      AS partition_name,
    pg_get_expr(child.relpartbound, child.oid) AS partition_bound
FROM pg_inherits AS i
JOIN pg_class      AS parent ON parent.oid = i.inhparent
JOIN pg_class      AS child  ON child.oid  = i.inhrelid
WHERE parent.relkind = 'p'
ORDER BY partition_name;
```

partition_name	partition_bound
orders_range_2022	FOR VALUES FROM ('2022-01-01') TO ('2023-01-01')
orders_range_2023	FOR VALUES FROM ('2023-01-01') TO ('2024-01-01')
orders_range_2024	FOR VALUES FROM ('2024-01-01') TO ('2025-01-01')
orders_list_customer_1	FOR VALUES IN (1)
orders_list_customer_2	FOR VALUES IN (2)
orders_list_customer_3	FOR VALUES IN (3)
orders_hash_customer_1	FOR VALUES WITH (modulus 4, remainder 0)
orders_hash_customer_2	FOR VALUES WITH (modulus 4, remainder 1)
orders_hash_customer_3	FOR VALUES WITH (modulus 4, remainder 2)
orders_hash_customer_4	FOR VALUES WITH (modulus 4, remainder 3)
orders_list_default	DEFAULT
orders_range_default	DEFAULT

Query - List only DEFAULT partitions

```
SELECT
    n.nspname          AS parent_schema,
    c.relname          AS parent_table,
    dp.oid::regclass  AS default_partition
FROM pg_partitioned_table AS pt
JOIN pg_class      AS c  ON c.oid = pt.partrelid
JOIN pg_namespace   AS n  ON n.oid = c.relnamespace
JOIN pg_class      AS dp ON dp.oid = pt.partdefid
WHERE pt.partdefid <> 0;

parent_schema | parent_table | default_partition
-----+-----+-----
public        | orders_range | orders_range_default
public        | orders_list  | orders_list_default
```

Query - List partitions and their sizes

```
SELECT
    t.relid::regclass                AS partition_name,
    t.level,
    pg_size.pretty(pg_relation_size(t.relid))      AS heap_size,
    pg_size.pretty(pg_total_relation_size(t.relid)) AS total_size_incl_indexes
FROM pg_partition_tree('public.orders_list'::regclass) AS t
WHERE t.isleaf
ORDER BY t.level, partition_name;

partition_name | level | heap_size | total_size_incl_indexes
-----+-----+-----+-----+
orders_list_customer_1 | 1 | 0 bytes | 8192 bytes
orders_list_customer_2 | 1 | 0 bytes | 8192 bytes
orders_list_customer_3 | 1 | 0 bytes | 8192 bytes
orders_list_default | 1 | 0 bytes | 8192 bytes
```

PostgreSQL Partitioning Settings

- Assume shared_buffers, work_mem, maintenance_work_mem, effective_cache_size are tuned
- Parameter **enable_partition_pruning**
 - Default value = ON
 - Skips scanning partitions that cannot match query predicates
 - For table inheritance, check constraints are required for pruning
- Parameter **constraint_exclusion**
 - Default value = 'partition'
 - Controls exclusion of partitions and tables based on constraints
 - For declarative partitioning, prefer enable_partition_pruning

- Additional parameters may impact partitioning performance and scalability
- Parameter **max_locks_per_transaction** - default 64
- Total locks = **max_locks_per_transaction** × (**max_connections** + **max_prepared_transactions**)
- High partition counts can exhaust lock slots, causing "out of shared memory" errors
- Parameter **max_pred_locks_per_relation** - default 64
- Sets the maximum number of predicate locks per relation
- Insufficient slots may also trigger "out of shared memory" errors

- Parameter **max_files_per_process**
 - Default value = 1000
 - Sets the maximum number of files a backend process can open simultaneously
 - High partition/table counts may require increasing this value
- Parameter **default_statistics_target**
 - Default value = 100
 - Determines the number of rows sampled for table statistics
 - Formula: $300 \times \text{default_statistics_target}$ rows sampled per column
 - The multiplier 300 is based on research for histogram accuracy (Ioannidis, 1998)
 - Higher values can improve query planning for partitioned tables
 - Larger targets increase ANALYZE time and statistics size

- Parameter **enable_partitionwise_aggregate**
 - Default value = OFF
 - Enables aggregation to run independently on each partition
 - Improves performance for GROUP BY and aggregate queries on partitions
- Parameter **enable_partitionwise_join**
 - Default value = OFF
 - Allows joins to be executed separately for matching partitions
 - Join condition must include all partitioning keys for partitionwise join
 - Reduces join costs for large partitioned tables with aligned partitioning
- Both settings may increase CPU and memory usage significantly
- Assess workload and resource impact before enabling in production

```
-- Create a partitioned table
CREATE TABLE sales (
    sale_id serial NOT NULL,
    sale_date date NOT NULL,
    amount numeric,
    PRIMARY KEY (sale_id, sale_date)
) PARTITION BY RANGE (sale_date);

-- Create partitions
CREATE TABLE sales_2023 PARTITION OF sales FOR VALUES FROM ('2023-01-01') TO ('2024-01-01');
CREATE TABLE sales_2024 PARTITION OF sales FOR VALUES FROM ('2024-01-01') TO ('2025-01-01');

-- Insert data into partitions
INSERT INTO sales_2023 (sale_date, amount) VALUES ('2023-05-01', 100);
INSERT INTO sales_2023 (sale_date, amount) VALUES ('2023-06-01', 150);
INSERT INTO sales_2024 (sale_date, amount) VALUES ('2024-02-01', 200);
INSERT INTO sales_2024 (sale_date, amount) VALUES ('2024-03-01', 250);
```

PostgreSQL Partitioning Settings - enable_partitionwise_aggregate



```
test=# set enable_partitionwise_aggregate = off;
SET
test=# explain SELECT sum(amount) AS total_sales FROM sales WHERE sale_date BETWEEN '2023-01-01' AND '2024-12-31';
          QUERY PLAN
```

```
Aggregate  (cost=56.09..56.10 rows=1 width=32)
->  Append  (cost=0.00..56.06 rows=12 width=32)
      ->  Seq Scan on sales_2023 sales_1  (cost=0.00..28.00 rows=6 width=32)
          Filter: ((sale_date >= '2023-01-01'::date) AND (sale_date <= '2024-12-31'::date))
      ->  Seq Scan on sales_2024 sales_2  (cost=0.00..28.00 rows=6 width=32)
          Filter: ((sale_date >= '2023-01-01'::date) AND (sale_date <= '2024-12-31'::date))
```

```
test=# set enable_partitionwise_aggregate = on;
SET
test=# explain SELECT sum(amount) AS total_sales FROM sales WHERE sale_date BETWEEN '2023-01-01' AND '2024-12-31';
          QUERY PLAN
```

```
Finalize Aggregate  (cost=56.08..56.09 rows=1 width=32)
->  Append  (cost=28.02..56.07 rows=2 width=32)
      ->  Partial Aggregate  (cost=28.02..28.03 rows=1 width=32)
          ->  Seq Scan on sales_2023 sales_1  (cost=0.00..28.00 rows=6 width=32)
              Filter: ((sale_date >= '2023-01-01'::date) AND (sale_date <= '2024-12-31'::date))
      ->  Partial Aggregate  (cost=28.02..28.03 rows=1 width=32)
          ->  Seq Scan on sales_2024 sales_1  (cost=0.00..28.00 rows=6 width=32)
              Filter: ((sale_date >= '2023-01-01'::date) AND (sale_date <= '2024-12-31'::date))
```

PostgreSQL Partitioning Settings - enable_partitionwise_join



```
-- Create a partitioned table for sales
CREATE TABLE sales (
    sale_id serial NOT NULL,
    sale_date date NOT NULL,
    amount numeric,
    PRIMARY KEY (sale_id, sale_date)
) PARTITION BY RANGE (sale_date);

-- Create partitions for sales
CREATE TABLE sales_2023 PARTITION OF sales FOR VALUES FROM ('2023-01-01') TO ('2024-01-01');
CREATE TABLE sales_2024 PARTITION OF sales FOR VALUES FROM ('2024-01-01') TO ('2025-01-01');

-- Create a partitioned table for customers
CREATE TABLE customers (
    customer_id serial NOT NULL,
    registration_date date NOT NULL,
    name text,
    PRIMARY KEY (customer_id, registration_date)
) PARTITION BY RANGE (registration_date);

-- Create partitions for customers
CREATE TABLE customers_2023 PARTITION OF customers FOR VALUES FROM ('2023-01-01') TO ('2024-01-01');
CREATE TABLE customers_2024 PARTITION OF customers FOR VALUES FROM ('2024-01-01') TO ('2025-01-01');

-- Insert data into sales partitions
INSERT INTO sales_2023 (sale_date, amount) VALUES ('2023-05-01', 100), ('2023-06-01', 150);
INSERT INTO sales_2024 (sale_date, amount) VALUES ('2024-02-01', 200), ('2024-03-01', 250);

-- Insert data into customers partitions
INSERT INTO customers_2023 (registration_date, name) VALUES ('2023-05-01', 'Alice'), ('2023-06-01', 'Bob');
INSERT INTO customers_2024 (registration_date, name) VALUES ('2024-02-01', 'Charlie'), ('2024-03-01', 'Dave');
```

```
test=# set enable_partitionwise_join = off;

test=# explain SELECT s.sale_date, s.amount, c.name FROM sales s
JOIN customers c ON s.sale_date = c.registration_date
WHERE s.sale_date BETWEEN '2023-01-01' AND '2024-12-31';

QUERY PLAN
-----
Hash Join  (cost=56.21..122.65 rows=144 width=68)
Hash Cond: (c.registration_date = s.sale_date)
-> Append  (cost=0.00..56.00 rows=2400 width=36)
    -> Seq Scan on customers_2023 c_1  (cost=0.00..22.00 rows=1200 width=36)
    -> Seq Scan on customers_2024 c_2  (cost=0.00..22.00 rows=1200 width=36)
-> Hash   (cost=56.06..56.06 rows=12 width=36)
    -> Append  (cost=0.00..56.06 rows=12 width=36)
        -> Seq Scan on sales_2023 s_1  (cost=0.00..28.00 rows=6 width=36)
            Filter: ((sale_date >= '2023-01-01'::date) AND (sale_date <= '2024-12-31'::date))
        -> Seq Scan on sales_2024 s_2  (cost=0.00..28.00 rows=6 width=36)
            Filter: ((sale_date >= '2023-01-01'::date) AND (sale_date <= '2024-12-31'::date))
```

```
test=# set enable_partitionwise_join = on;

test=# explain SELECT s.sale_date, s.amount, c.name FROM sales s
JOIN customers c ON s.sale_date = c.registration_date
WHERE s.sale_date BETWEEN '2023-01-01' AND '2024-12-31';

QUERY PLAN
-----
Append  (cost=28.07..110.23 rows=72 width=68)
-> Hash Join  (cost=28.07..54.94 rows=36 width=68)
    Hash Cond: (c_1.registration_date = s_1.sale_date)
        -> Seq Scan on customers_2023 c_1  (cost=0.00..22.00 rows=1200 width=36)
        -> Hash  (cost=28.00..28.00 rows=6 width=36)
            -> Seq Scan on sales_2023 s_1  (cost=0.00..28.00 rows=6 width=36)
                Filter: ((sale_date >= '2023-01-01'::date) AND (sale_date <= '2024-12-31'::date))
-> Hash Join  (cost=28.07..54.94 rows=36 width=68)
    Hash Cond: (c_2.registration_date = s_2.sale_date)
        -> Seq Scan on customers_2024 c_2  (cost=0.00..22.00 rows=1200 width=36)
        -> Hash  (cost=28.00..28.00 rows=6 width=36)
            -> Seq Scan on sales_2024 s_2  (cost=0.00..28.00 rows=6 width=36)
                Filter: ((sale_date >= '2023-01-01'::date) AND (sale_date <= '2024-12-31'::date))
```

PostgreSQL Extensions for Partitioning

- pg_partman - Advanced partition management, supports time and integer ranges
- TimescaleDB - Time-series extension, automatic partitioning via hypertables
- pg_repack - Online table and index reorganization, enables clustering
- pg_squeeze - Automates bloat removal, supports online table clustering
- Citus - Distributed extension, enables sharding and multi-node partitioning

- Does not support partitioning
- But can be life saver in case of badly bloated tables/indexes
- Extension for online table and index reorganization
- Removes bloat from tables and indexes without locks
- Supports clustering tables based on an index
- Creates a new copy of the table with desired order
- Swaps the new table in place with minimal locking
- Requires minimal downtime during reorganization
- pg_squeeze is designed to replace pg_repack

- Automates creation and management of time- and number-based partitions
- Supports native declarative partitioning; inheritance deprecated
- Includes background worker for automatic partition maintenance
- Configure with `shared_preload_libraries = 'pg_partman_bgw'`
- Eliminates need for external schedulers like cron
- Requires PostgreSQL 14 or newer
- Recommended by AWS RDS for scalable partition management

- Extension for scalable time-series data on PostgreSQL
- Time-series tables are managed as "hypertables"
- Hypertables are partitioned into chunks automatically
- Supports partitioning by time, or by time and space (e.g., device_id)
- Enables continuous aggregations via specialized materialized views
- Optimized for high ingest rates and analytical queries
- Provides native compression for historical data



TIMESCALE

Thank you for your attention!



All my slides



Recorded talks

