

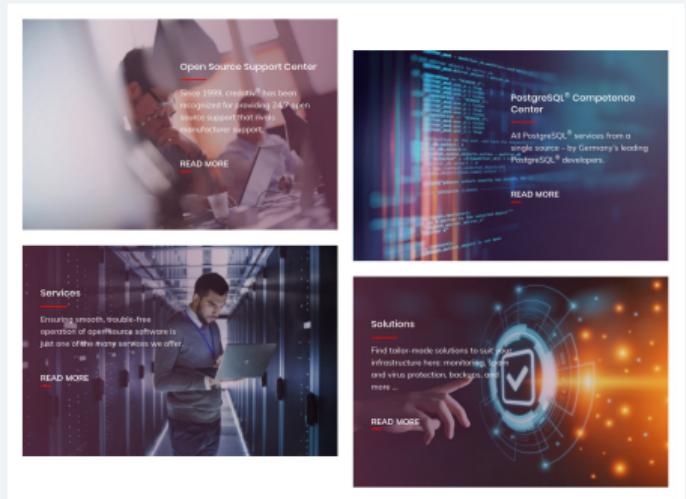
PostgreSQL 18 Asynchronous Disk I/O

How it works under the hood

Josef Machytka <josef.machytka@credativ.de>

2025-10-24 - datativ Tech Talk

- Founded 1999 in Jülich, Germany
- Close ties to Open-Source Community
- More than 40 Open-Source experts
- Consulting, development, training, support (3rd-level / 24x7)
- Open-Source infrastructure with Linux, Kubernetes, Proxmox
- Open-Source databases with PostgreSQL
- DevSecOps with Ansible, Puppet, Terraform and others
- Since 2025 independent owner-managed company again



- Professional Service Consultant - PostgreSQL specialist at credativ GmbH
- 33+ years of experience with different databases
- PostgreSQL (13y), BigQuery (7y), Oracle (15y), MySQL (12y), Elasticsearch (5y), MS SQL (5y)
- 10+ years of experience with Data Ingestion pipelines, Data Analysis, Data Lake and Data Warehouse
- 3+ years of practical experience with different LLMs / AI / ML including architecture and principles
- From Czechia, living now 12 years in Berlin

-  linkedin.com/in/josef-machytka
-  medium.com/@josef.machytka
-  youtube.com/@JosefMachytka
-  github.com/josmac69/conferences_slides
-  researchgate.net/profile/Josef-Machytka

All My Slides:

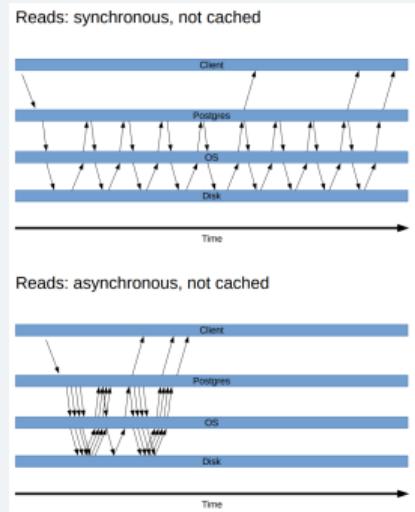


Recorded talks:



Streaming I/O in PostgreSQL 17

- PG17 introduced "streaming I/O" - multiple pages in one request
 - More systematically issues POSIX_FADV_WILLNEED for random access
 - Key posix_fadvise hint to OS for better prefetching
 - Asks kernel to prefetch multiple pages in advance
- Vectored I/O requests using readv/writev system calls
- Perform multiple page reads/writes in one syscall
- Improves sequential scan performance by up to 30%
- But still synchronous, one-at-a-time per backend process
- Does not utilize full capabilities of modern storage devices
- Waiting for Postgres 17: Streaming I/O for sequential scans & ANALYZE



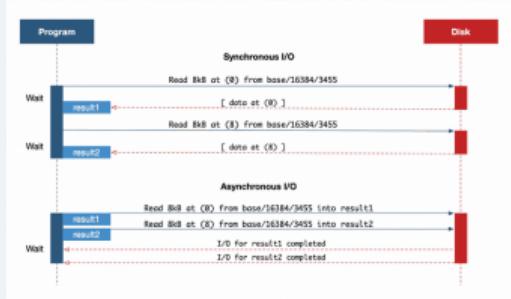
Images from the talk

The path to using AIO in Postgres (PGConf.EU 2023)

Asynchronous I/O

- New AIO subsystem allows multiple concurrent file reads
- Current implementation is only minimal
- Does not support writes - OLTP operations do not benefit
- Multi process model made implementation challenging

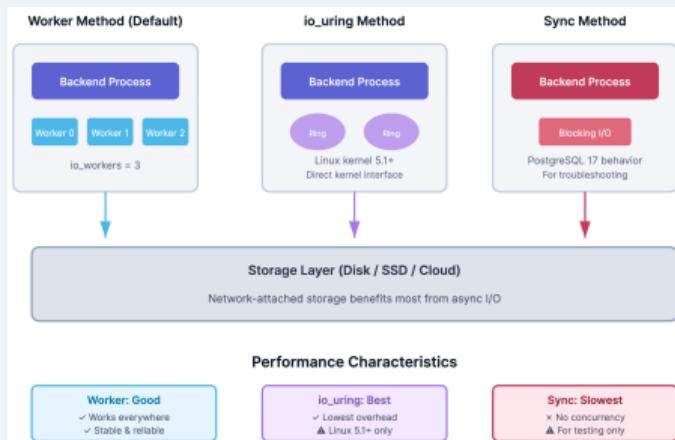
- PG18 Async I/O significantly improves performance of:
 - Large sequential scans, GROUP BY, COPY (up to 2-4x faster)
 - Bitmap heap scans
 - Vacuum (some benchmarks up to 3-4x faster, some 0)
 - Read operations on cloud DBs with network-attached storage



Images from the article
[Accelerating Disk Reads with Asynchronous I/O](#)

Async I/O: Configuration Options

- **io_method:**
 - worker - default, pool of I/O worker processes
 - io_uring - Linux-specific async I/O queues
 - sync - "fallback" synchronous I/O, but still uses AIO API
- **io_workers:**
 - 3 - default, too low for larger systems
 - Start with 1/4 of total CPU threads
 - Some benchmarks suggest up to 1/2 of threads

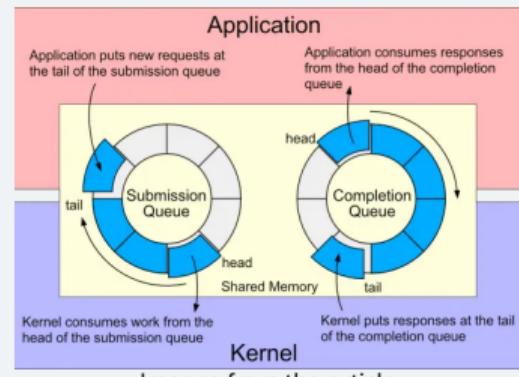


Images from the article

[PostgreSQL 18 Asynchronous I/O: A Complete Guide](#)

Linux io_uring method - overview

- Modern Linux I/O interface - needs Linux 5.1+
 - check with `uname -r` command -> shows kernel version
- Requires Postgres build with `--with-liburing` / `-Dliburing`
 - check with `pg_config --configure | grep --color=always liburing`
 - select `pg_config` from `pg_config()` where `pg_config:text ilike '%liburing%'`;
 - -> both show if liburing is enabled in build
 - `ldd /path/to/postgres | grep uring` -> shows if liburing is linked
- Backends submit I/O requests directly into ring buffer
- Kernel processes requests asynchronously
- Completion queue entries consumed by backends
- Fewer context switches and wakeups
- Better latency on fast NVMe and network-attached SSDs



Images from the article
[Why you should use io_uring](#)

Linux io_uring method – Security Risks



- io_uring bypasses traditional I/O system read/write calls
 - Including network sends/receives
 - Rootkits and malware can exploit it to hide I/O activity
 - Google reported 60% of kernel exploits in 2022 used io_uring
 - Some security tool do not detect io_uring activity
 - Sysdig added new detection rules for io_uring abuse
 - Therefore some container runtimes disable it for security
 - io_uring also complexity led to numerous bugs
-
- ARMO: [io_uring Is Back, This Time as a Rootkit](#)
 - Sysdig: [Detecting and Mitigating io_uring Abuse for Malware Evasion](#)
 - [io_uring: Linux Performance Boost or Security Headache?](#)

A screenshot of the Sysdig dashboard. At the top, a banner reads "Suspicious io_uring Activity Detected". Below it, a table shows a single row of data:

Severity	Medium
Policy	Syndy Runtime Violable Events
Rule name	Suspicious io_uring Activity Detected
Rule type	Failure / Syndy

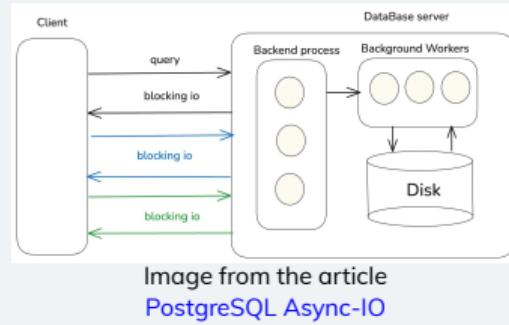
Under the "Process" section, there is a table with the following data:

Name	client
Path	/sys/fs/bus/usb/devices/1d:11/uevent
Parent name	bus
Executable path	/home/vagrant/moving/test/client
Executable file	client
Name and arguments	client
PID	2019
Parent name and arguments	bus
Parent PID	2036

Images from the Sysdig article

Async io_method = worker: Overview

- Uses dedicated background I/O worker processes
- Backends submit I/O requests to shared memory queue
- I/O workers pick up requests and perform them asynchronously
- Completion queue entries consumed by backends
- More context switches and wakeups than io_uring
- But works on all OSes that support POSIX AIO API
- Recommended for non-Linux systems or older Linux kernels
- Some benchmarks show better performance to io_uring
- Avoids security risks of io_uring method



- Monitored via new pg_aios view:
 - Shows current I/O requests - status, result etc.
 - Status: DEFINED, STAGED, SUBMITTED, COMPLETED_IO, ...
 - Result: UNKNOWN, OK, PARTIAL, WARNING, ERROR
- Enhanced pg_stat_io view - new columns: read_bytes, write_bytes, extend_bytes
- Parameter track_io_timing
 - Enables additional I/O timing statistics
 - Default is off - if on, adds some overhead
- On Linux level using iotop command

```
-- output from iotop command:
```

TID	PRI	USER	DISK READ	DISK WRITE>	COMMAND
617931	be/4	postgres	32.17 M/s	0.00 B/s	postgres: io worker 0
617932	be/4	postgres	72.11 M/s	0.00 B/s	postgres: io worker 1
617933	be/4	postgres	41.50 M/s	0.00 B/s	postgres: io worker 2
617934	be/4	postgres	58.00 M/s	0.00 B/s	postgres: io worker 4

- **io_combine_limit:**
 - 16 - default (i.e. 16 pages = 128kb) - dynamic
 - Without unit - number of pages, with kB/MB - size in kB/MB
 - How many pages can be combined in 1 request
 - Optimal value depends on underlying HW / OS capabilities
 - Hardware and protocols have segment and size limits
 - After some threshold, increasing value has no effect
 - Limited internally in PG by setting io_max_combine_limit
- **io_max_combine_limit:**
 - 16 - default (i.e. 16 pages = 128kb) - requires restart
 - Limits io_combine_limit - maximum pages per request
 - Cluster wide limit for all backends
 - Typical max: Unix 128 (1 MB), Windows 16 (128 kB)
 - Added to manage memory consumption & structure sizes

- **io_max_concurrency:**

- maximum number of concurrent I/O operations per process (capped to 64)
- default -1 = selects number based on shared_buffers and max number of backends
- Requires restart to change

- **effective_io_concurrency:**

- number of concurrent I/O operations that can be executed simultaneously
- new default 16, range 1-1000, value 0 disables async requests
- Higher values for faster storage - 200 for NVMe SSDs
- Very high values may increase I/O latency for all queries
- Can be changed dynamically

How it works under the hood

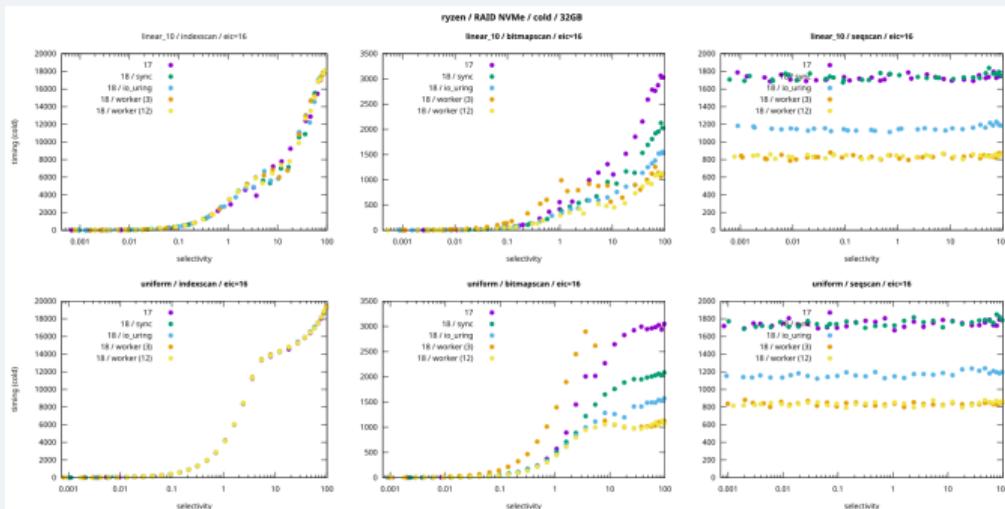
- `src/backend/storage/aio/README.md`
- [Asynchronous & Direct IO \(PostgreSQL Source Code git master\)](#)
- `src/backend/storage/aio/method_worker.c`
- Postmaster creates on startup a pool of background worker processes
- Each worker is a separate OS process
- Workers still perform synchronous IO, they just read in parallel
- Blocking syscalls are moved to pool of separate processes
- Workers ignore SIGTERM termination signal
- Get explicit shutdown via SIGUSR2 later in the shutdown sequence
- If worker fails, it marks IO request as failed and exists
- Postmaster starts a new worker to replace it

How it works under the hood

- `src/backend/storage/aio/method_worker.c`
- Each worker loops forever, waiting for requests
- Request for I/O sent to shared memory submission queue
- Worker wakes up, reads requests from shared memory queue
- Each woken IO worker can wake 2 more - note about future improvement to N
- Workers perform synchronous read syscalls for each request
- Put results into shared memory completion queue
- Notify backend process via latch

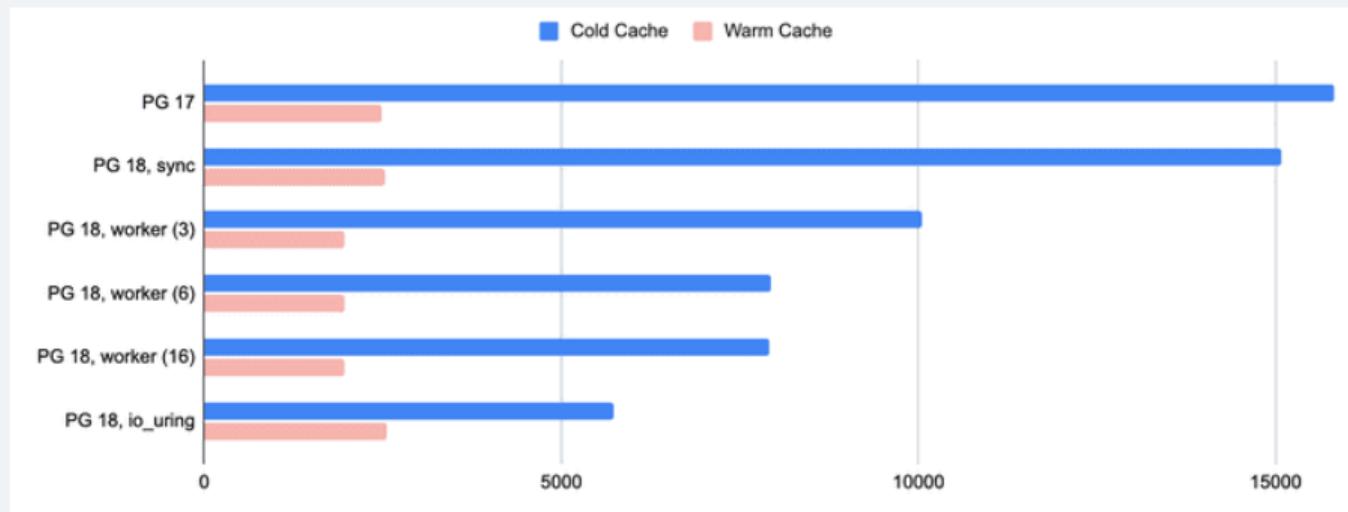
Async I/O Benchmarks

- Tomas Vondra: Tuning AIO in PostgreSQL 18:
- Benchmarks Ryzen 9900X - 12 cores/ 24 threads, 32 GB RAM, 4 NVMe SSD RAID0
- Tests for PG 17, PG 18 with sync, PG 18 with io_uring, PG 18 with AIO worker (3 and 12)
- (uniform = entirely random distribution, linear_10 = sequential with 10% randomness)
- index scan - no effect, index scans do not use AIO yet, only sync I/O
- bitmap scan - 3 workers slower than sync for low selectivity queries, 12 workers 2x faster
- sequential scan - workers 2x faster than sync, io_uring in between



Async I/O Benchmarks

- Waiting for Postgres 18: Accelerating Disk Reads with Asynchronous I/O:
- AWS c7i.8xlarge instance 32 vCPUs, 64 GB RAM, 100GB io2 EBS volume 20,000 IOPS



- Major core operations that benefit from AIO:
 - Heap **sequential scans** - plain SELECT and COPY style scans
 - **ANALYZE** sampling on heap tables
 - **VACUUM** of heap tables and B-tree indexes
 - **Bitmap heap scans** when the executor decides to use them
- These operations usually know many future block numbers in advance
- Can keep multiple reads in flight while doing useful CPU work
- Operations that do not yet use AIO:
 - B-tree index scans (including index-only scans)
 - Recovery and replication code paths
 - Write operations (INSERT, UPDATE, DELETE, etc.)
 - Small OLTP lookups -> index + single heap page

- Autovacuum workers use the same VACUUM and ANALYZE code paths
- Large tables maintained by autovacuum can also profit from async reads
- AIO is intentionally focused on large, predictable scans
 - Avoids fighting the kernel's own readahead heuristics
 - Lets PostgreSQL drive readahead based on query plans
- Future work will add more operations into read streams and AIO
- Definitely index-only scans and eventually writes

- Asynchronous & Direct IO (PostgreSQL Source Code git master)
- Tomas Vondra: Tuning AIO in PostgreSQL 18
- Waiting for Postgres 18: Accelerating Disk Reads with Asynchronous I/O
- Boosting Typical Query Patterns - PostgreSQL 18's Performance Enhancements (PGConf.EU 2025)
- Get Excited About Postgres 18
- What went wrong with AIO (PGconf.DEV 2025)
- PostgreSQL 18 Asynchronous I/O (Neon blog)
- PostgreSQL 18: A Comprehensive Guide to New Features for DBAs and Developers

Thank you for your attention!



All my slides

