

Fast & Prepared

Prepared Statements in PostgreSQL

Josef Machytka <josef.machytka@credativ.de>

2025-11-14 - datativ Tech Talk

- Founded 1999 in Jülich, Germany
- Close ties to Open-Source Community
- More than 40 Open-Source experts
- Consulting, development, training, support (3rd-level / 24x7)
- Open-Source infrastructure with Linux, Kubernetes, Proxmox
- Open-Source databases with PostgreSQL
- DevSecOps with Ansible, Puppet, Terraform and others
- Since 2025 independent owner-managed company again



- Professional Service Consultant - PostgreSQL specialist at credativ GmbH
- 33+ years of experience with different databases
- PostgreSQL (13y), BigQuery (7y), Oracle (15y), MySQL (12y), Elasticsearch (5y), MS SQL (5y)
- 10+ years of experience with Data Ingestion pipelines, Data Analysis, Data Lake and Data Warehouse
- 3+ years of practical experience with different LLMs / AI / ML including architecture and principles
- From Czechia, living now 12 years in Berlin

- **LinkedIn:** linkedin.com/in/josef-machytka
- **Medium:** medium.com/@josef.machytka
- **YouTube:** youtube.com/@JosefMachytka
- **GitHub:** github.com/josmac69/conferences_slides
- **ResearchGate:** researchgate.net/profile/Josef-Machytka
- **Academia.edu:** academia.edu/JosefMachytka
- **Sessionize:** sessionize.com/josefmachytka

All My Slides:



Recorded talks:



Topics Covered

- Why, what, when to use, when to avoid
- Protocol-level vs SQL commands
- EXPLAIN, plan caching, generic vs custom plans
- Prepared statements & partitioned tables
- Internal implementation overview
- Plan invalidation & replanning
- Hidden usage in FDWs, PL/pgSQL
- Common use cases

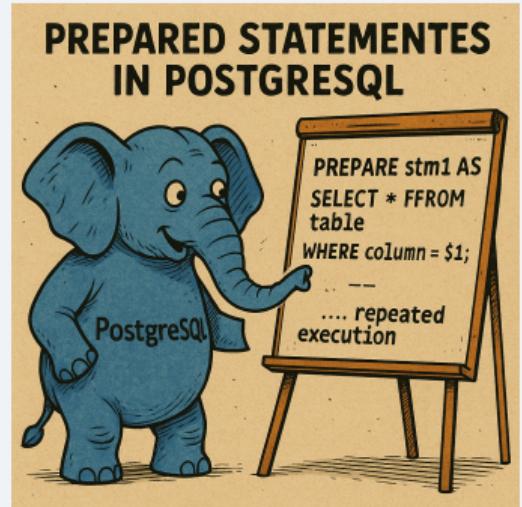


Image credit ChatGPT

Why Prepared Statements?

- What does postgres do with a query?
 - Parses the SQL statement
 - Analyzes it
 - Rewrites it
 - Plans it
 - Executes the plan & Returns the result
- Every step increases total query runtime
- Prepared statements can optimize this process
 - Parse, Analyze, Rewrite, Plan only once
 - Execute many times

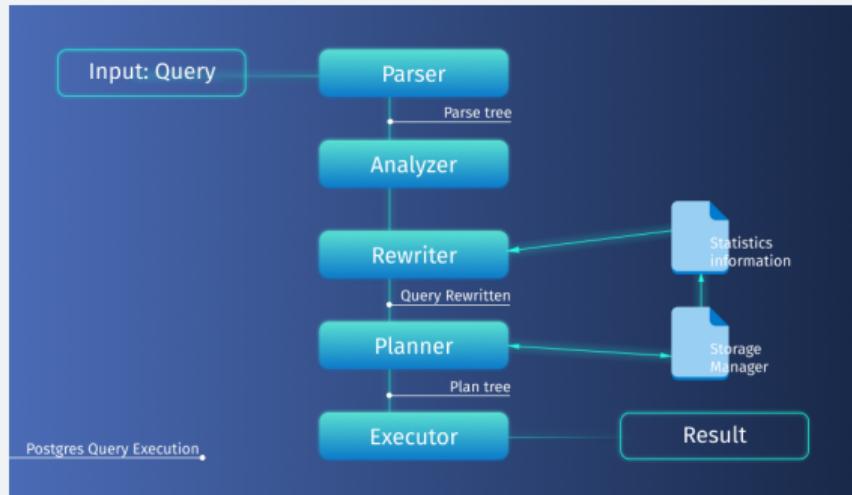


Image from the article
[PostgreSQL EXPLAIN for Analyzing Query Execution Plans](#)

What Are These Prepared Statements Anyway?

- Prepared statements are stored in a server-side cache
- They are kept in memory, session-scoped
- Survive across transactions & multiple queries
- **Server-side** prepared statements:
 - - managed by PostgreSQL server -
 - - created using SQL commands PREPARE, EXECUTE, DEALLOCATE
- **Client-side** prepared statements:
 - - managed by client libraries
 - - emulated inside library, not created on server

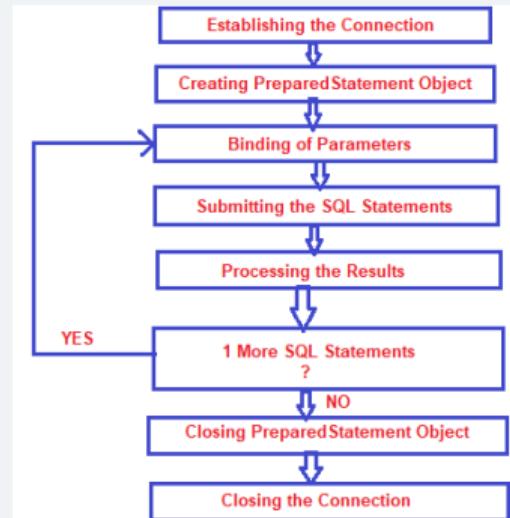


Image from the article
[Prepared Statement in JDBC](#)

Prepared Statements vs Prepared Transactions?

- **Prepared Statements** are cached query plans
- They are performance optimization feature
- **Prepared Transactions** are 2-phase commit transactions
- Prepared Transactions are managed with:
 - PREPARE TRANSACTION;
 - COMMIT PREPARED;
 - ROLLBACK PREPARED;
- They are internal PostgreSQL feature for distributed transactions
- Intended for ensuring atomicity across multiple databases
- Not related to Prepared Statements

When Shall We Use Prepared Statements?

- When a single session executes the same query multiple times
- Just with different parameters
- When performance is critical
- And query planning overhead needs to be minimized
- Advantages of using prepared statements:
 - Prevent repeated parsing & planning overhead
 - Prevent SQL injection attacks
 - Reduce network traffic
 - Improve code readability

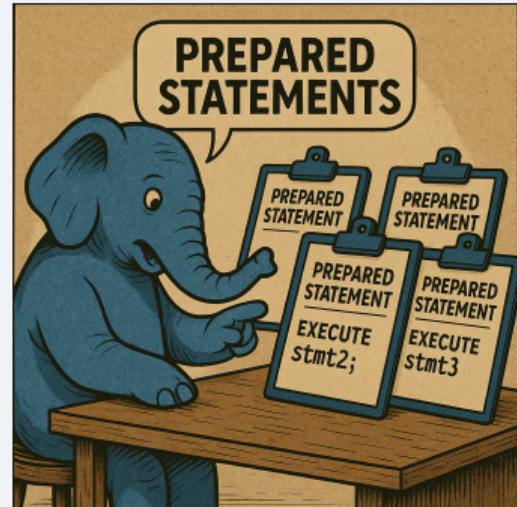


Image credit: ChatGPT

When Shall We NOT Use Prepared Statements?

- When queries are executed only once
- When queries are highly dynamic
- When query planning time is negligible
- When using connection pooling
- When prepared statement cache size is limited
- When connection lifetime is short
- Gotchas:
 - Can switch to bad plans over time
 - After 5th execution, re-planning happens (default)
 - Generic plans can be much less efficient
 - Increased memory usage with many prepared statements

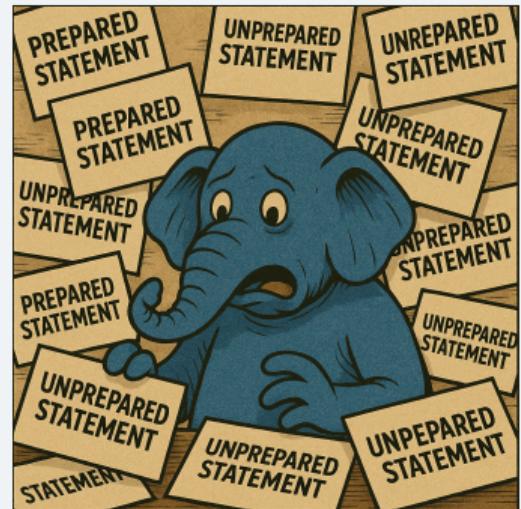


Image credit: ChatGPT

- **SQL-level commands** - standard SQL commands:
 - PREPARE, EXECUTE, DEALLOCATE
 - Must be parsed as text to be understood
- **Protocol-level commands** - "wire protocol" between client & server
 - Parse - create a prepared statement on the server
 - Bind - attach parameters to a prepared statement
 - Execute - execute the prepared statement with parameters
 - Close - deallocate the prepared statement (added in PostgreSQL 17)
- See in PostgreSQL documentation [Chapter 54. Frontend/Backend Protocol](#)
- Part [54.2.3. Extended Query](#) describes prepared statements
- Most client libraries use protocol-level commands
- PgBouncer can intercept protocol-level commands & handle them in transaction mode

Prepared Statements and Connection Pooling

- Prepared statements are session-scoped
- Some poolers do not support them in transaction mode
- PgBouncer supports protocol-level prepared statements
- But not SQL-level prepared statements
- Feature implemented in versions 1.21 and 1.22
- Functionality requires PostgreSQL 17 or higher

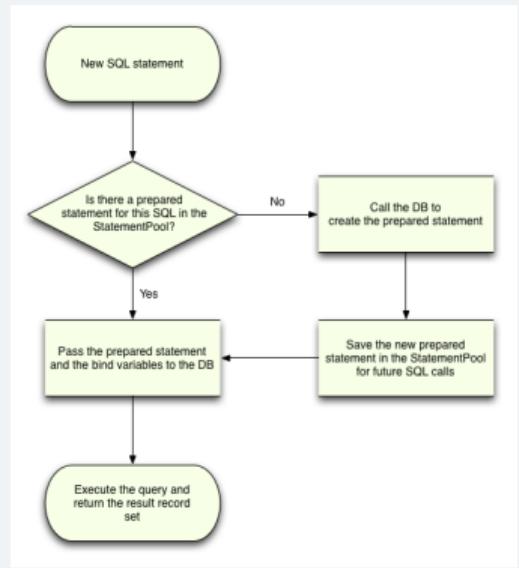


Image from the article
[Show some love for prepared statements in Rails 3.1](#)

View: pg_prepared_statements

- System view showing prepared statements in current session
- Columns:
 - name - name of the prepared statement
 - statement - original query string
 - prepare_time - timestamp of when it was prepared
 - parameter_types - array of parameter data types
 - result_types - array of result column data types
 - from_sql - boolean indicating if prepared via SQL commands (true) or protocol (false)
 - generic_plans - number of times generic plan was used
 - custom_plans - number of times custom plan was used
- Useful for monitoring prepared statements usage

EXPLAIN command & Prepared Statements

- EXPLAIN shows the execution plan of a query for specific parameters
- We can test prepared statements with new GENERIC_PLAN option

```
CREATE TABLE test (id int PRIMARY KEY, data text);
-- Insert data, run ANALYZE

EXPLAIN GENERIC_PLAN SELECT * FROM test WHERE id = \$1;

PREPARE my_stmt (int) AS SELECT * FROM test WHERE id = \$1;

SET plan_cache_mode = 'force_generic_plan';
EXPLAIN EXECUTE my_stmt (42);

SET plan_cache_mode = 'force_custom_plan';
EXPLAIN EXECUTE my_stmt (42);
```

- `plan_cache_mode` (default: `auto`)
 - Controls the behavior of the plan cache for prepared statements
 - Options: `auto`, `force_custom_plan`, `force_generic_plan`
 - `auto` = plan is reevaluated after 5 executions
- First 5 executions use custom plans (`plan_cache_mode = auto`)
- On 6th execution, PostgreSQL compares generic vs average custom plan costs

- **Generic plan**
 - Parameter agnostic - "one size fits all"
 - Uses average distribution of parameter values
 - After creation does not require replanning (no planning costs)
 - Generic plan cannot prune partitions - validates/locks all partitions
- **Custom plan**
 - Optimized for specific parameter values based on statistics
 - Code includes costs of planning into custom plan costs
- Index scan vs sequential scan flip is common difference
- PostgreSQL switches to generic plan when:
 - Generic plan costs are cheaper than execution + replanning costs of custom plans
- For table with many partitions force custom plan

- Partitioned tables have significant difference between generic & custom plans
- Custom plans can prune partitions based on parameter values
- During PREPARE/Parse backend will:
 - Parse-analyze the query
 - Create parse tree in a cached plan structure
 - But no parameter values are known yet
- During EXECUTE/Bind:
 - It will create the actual plan
 - With custom plan partitions can be pruned now
- Generic plans must validate/lock all partitions
- This can lead to significant overhead with many partitions
- Recommendation: force custom plans for partitioned tables
- Postgres: on performance of prepared statements with partitioning

What PostgreSQL Does Internally

- When PREPARE or protocol Parse command is received:
 - Name is validated (not empty, not duplicate)
 - Internal memory structure are prepared - raw parse tree etc.
 - Parameters types are determined and converted into OIDs
 - Query is parsed, analyzed, and rewritten
 - Plan is completed and stored in plansource structure
 - All metadata is stored in PreparedStatement hash table

What PostgreSQL Does Internally

- When EXECUTE or protocol Execute command is received:
 - PreparedStatement hash table is looked up by name
 - Parameters are evaluated and converted to internal format
 - In case of partitioned table, pruned plan is created now
 - Process creates "portal" for execution - query instance with parameters
 - Cached plan is obtained from plansource structure
 - Query "portal" is associated with the plan
 - Portal is executed, results returned, portal deallocated

- `src/backend/commands/prepare.c`
- Implements SQL commands PREPARE, EXECUTE, DEALLOCATE
- As well as memory structures for prepared statements
- Defines per-backend hash table to store prepared statements
- Structure `PreparedStatement` holds:
 - Name of the prepared statement
 - Cached query plan (plansource structure - all details of the plan, quite big)
 - Boolean - prepared via SQL (true) or protocol (false)
 - Timestamp of creation

- Plan cache management is in `src/backend/utils/cache/plancache.c`
 - Decides when to use generic vs custom plans
 - Tracks schema changes invalidating cached plans
 - Defines plansource structure as well
- Prepared statements memory structures - use several memory contexts:
 - Source context - holds CachedPlanSource and parse tree
 - Query context - holds analyzed and rewritten query tree & dependency data
 - Plan context - for each CachedPlan, created by BuildCachedPlan
- Memory contexts are cleaned up when prepared statement is deallocated

Implementation of protocol-level Prepared Statements

- Implemented in the libpq client library
- Functions:
 - `PQprepare()` - prepares a statement - sends Parse command to server
 - `PQexecPrepared()` - executes a prepared statement - sends Bind & Execute
 - `PQdescribePrepared()` - sends Describe command to get metadata about prepared statement
 - `PQclear()` - deallocates a prepared statement - sends Close command to server (PostgreSQL 17+)
- Connection libraries use these functions to manage prepared statements
- Described in PostgreSQL documentation [32.3. Command Execution Functions](#)

- Cached plans can become invalid due to schema changes
- Code in `plancache.c` registers "invalidation callbacks" on system catalog
- Processed by `src/backend/storage/ipc/sinval.c` code -> "sinval" events
- When an invalidation event matches a cached plan's dependencies, it is marked as invalid
- During next execution, invalid plan is re-analyzed or re-planned
- Usage of generic vs custom plan is re-evaluated based on current parameters
- New statistics are stored in the plansource structure
- New numbers are reported in `pg_prepared_statements` view

Hidden Usage of Prepared Statements

- Some PostgreSQL features use prepared statements internally
- Example is `postgres_fdw` foreign data wrapper
- See in `contrib/postgres_fdw/postgres_fdw.c`
- It creates prepared statements for execution of INSERT/UPDATE/DELETE
- This optimizes performance when accessing foreign tables

- PL/pgSQL code automatically turns static queries into prepared statements
- This improves performance of repeated executions of functions
- Dynamic statements must be parsed each time, cannot be prepared

- Web applications with high frequency of SELECT queries
 - Same SELECTs with different parameters
 - Reduces parsing/planning overhead significantly (up to 36% in some benchmarks)
- Data ingestion pipelines / ETL with repeated INSERT/UPDATE operations
 - Prepared statements speed up data loading by reusing plans
 - Protect against SQL injection attacks
- Repeated complex analytical queries in reporting applications
 - Complicated queries often have high planning overhead
 - Planning often takes longer than execution

Additional Resources



- Articles:
 - [PostgreSQL Performance Insights: Extended Query Protocol vs. Simple Query Protocol](#)
 - [EXPLAIN \(GENERIC_PLAN\): New in PostgreSQL 16](#)
 - [Puzzling Postgres: solving an unreproducible performance issue](#)
 - [Supporting Postgres Named Prepared Statements in Hyperdrive](#)
 - [When are PostgreSQL prepared statements fully prepared for execution?](#)
 - [More Memory, More Problems \(generic plans incident\)](#)
- Talks & Videos:
 - [Comparing Postgres connection pooler support for prepared statements | POSETTE 2024](#)

Thank you for your attention!



All my slides



Recorded talks

