

# The Alchemy of Shared Buffers

Balancing Concurrency and Performance

---

Josef Machytka <josef.machytka@credativ.de>

2026-02-06 - CERN PGDay 2026

- Founded 1999 in Jülich, Germany
- Close ties to Open-Source Community
- More than 40 Open-Source experts
- Consulting, development, training, support (3rd-level / 24x7)
- Open-Source infrastructure with Linux, Kubernetes, Proxmox
- Open-Source databases with PostgreSQL
- DevSecOps with Ansible, Puppet, Terraform and others
- Since 2025 independent owner-managed company again



- Professional Service Consultant - PostgreSQL specialist at credativ GmbH
- 33+ years of experience with different databases
- PostgreSQL (13y), BigQuery (7y), Oracle (15y), MySQL (12y), Elasticsearch (5y), MS SQL (5y)
- 10+ years of experience with Data Ingestion pipelines, Data Analysis, Data Lake and Data Warehouse
- 3+ years of practical experience with different LLMs / AI / ML including architecture and principles
- From Czechia, living now 12 years in Berlin

-  [linkedin.com/in/josef-machytka](https://linkedin.com/in/josef-machytka)
-  [medium.com/@josef.machytka](https://medium.com/@josef.machytka)
-  [youtube.com/@JosefMachytka](https://youtube.com/@JosefMachytka)
-  [github.com/josmac69/conferences\\_slides](https://github.com/josmac69/conferences_slides)
-  [researchgate.net/profile/Josef-Machytka](https://researchgate.net/profile/Josef-Machytka)

All My Slides:




Recorded talks:



# PostgreSQL Multi-Process Architecture

- PostgreSQL uses a multi-process architecture
  - Operates as a collection of cooperating processes
  - Every connection, every background task is a separate OS process
  - Processes communicate via shared memory inter-process communication (IPC)



- Linux philosophy: Everything is a file -> dentry (directory entry) & inode structures
- Shared memory on Linux implemented via `tmpfs` filesystem
- It is a filesystem interface to access memory as files
- Introduced in Linux kernel 2.4 (2001) as successor of older `ramfs`
- Internally used even if `tmpfs` is disabled for users
- `Tmpfs` can be used by SysV IPC (`shmget`, `shmat`) and `mmap` interfaces
- PostgreSQL originally used SysV shared memory segments
- Since PG 9.3 POSIX shared memory via `mmap` is default on Linux




Image from the article [Shared memory on Linux](#)

- Users usually interact with tmpfs via `/dev/shm` directory
- This filesystem grows and shrinks dynamically as files are created or deleted
- Size is limited by available RAM and swap space
- Uses page cache - file I/O is done directly in memory
- Writing allocates physical memory pages - associated with `dentry` and `inode` structures
- PostgreSQL uses `/dev/shm` for communication between processes
- Has small default on docker (64MB), can be exhausted quickly
- Should be increased with `--shm-size` docker parameter or `shm_size` in docker-compose
- If `/dev/shm` is exhausted, PG reports "could not resize shared memory segment" error

```
my-linux ~ $ df -h | grep tmpfs
tmpfs              3.2G  3.7M  3.1G  1% /run
tmpfs              16G  543M  15G  4% /dev/shm
tmpfs              5.0M  8.0K  5.0M  1% /run/lock
tmpfs              3.2G  140K  3.2G  1% /run/user/1000
```

- Data pages in `tmpfs` are anonymous memory pages
  - -> not backed by any file on disk, contents exist only in RAM
- Can be swapped out if necessary - like other memory pages
  - Kernel's page-replacement algorithm decides when
  - Can be security risk for sensitive data - encrypted swap recommended
  - -> swapping can degrade PostgreSQL performance
- `tmpfs` supports Transparent Huge Pages (THP)
  - Improves performance for large memory allocations
  - But can cause performance degradation for PostgreSQL
  - -> PG docs recommend to disable THP for DB servers
- It has also NUMA allocation policy option
  - -> local on current CPU / bind to specific NUMA node
- Oversizing `tmpfs` & swap disabled can deadlock the system
  - -> if OOM killer is disabled / cannot free mem

- Originally PostgreSQL used `System V` (SysV) for Inter-Process Communication (IPC)
  - Usage of SysV shared memory segment was later discouraged
  - Its default size was limited by default to 32MB
  - Higher sizes required reconfiguration of OS kernel parameters
  - See in [Absurd Shared Memory Limits](#) blog post by Robert Haas (2012)
  - Documentation still shows [how to configure SysV shared memory](#)
  - SysV allowed PG to detect multiple postmasters accessing the same data directory
  - Now this is done via `postmaster.pid` file -> empty PID file can prevent start (!)
- Since PG 9.3, default is POSIX shared memory on Linux
- Parameter `shared_memory_type` controls the type of shared memory
  - `mmap` - for anonymous shared memory allocated using mmap (default on Linux)
  - `sysv` - for System V shared memory allocated via shmget
  - `windows` - for Windows named shared memory

# How PostgreSQL Shares Memory Between Processes

---

- Main shared memory area allocated on server startup by the `postmaster`
- Postmaster sits in loop waiting for connections -> woken by kernel -> accepts connection
- Creates new process for a connection using `fork()` system call
  - Child process inherits parent's memory mapping - mapped as `MAP_SHARED`
  - All processes see the same shared memory region
  - Changes in memory are visible to all processes
- PostgreSQL can use `Huge Pages` for shared buffers and some other shared memory objects
  - Reduce TLB (Translation Lookaside Buffer) misses
  - Hence reduce CPU usage = improve performance

# How To Properly Use Huge Pages in PostgreSQL

---

- Usage of Huge Pages is not entirely straightforward
  - By default Huge Pages are not enabled on Linux -> `vm.nr_hugepages = 0`
  - Check `sysctl vm.nr_hugepages` to see the current setting
  - -> PG setting `huge_pages = try` does not have any effect
- Wrongly configured Huge Pages can cause memory issues
  - Configured number of Huge Pages is pre-allocated at Linux boot time
  - Pages are pinned in memory -> not swappable, guaranteed TLB efficiency
  - Only processes which implement usage of Huge Pages can use this area
  - It is not available for other processes
  - Configured number can be dynamically changed at runtime
  - But only downsizing is recommended

# How To Properly Use Huge Pages in PostgreSQL

---

- `cat /proc/meminfo |grep -i hugepagesize` -> see Huge Page size
- Typical Huge Page size is 2 MB, but 1 GB is also possible
- -> read-only parameter `shared_memory_size_in_huge_pages`
- -> shows how much huge pages would be required
- -> can be checked before starting PostgreSQL -  
`postgres -D $PGDATA -C shared_memory_size_in_huge_pages`
- Example: 8GB shared buffers -> 4096 Huge Pages 2048kB big
- But check of `shared_memory_size_in_huge_pages` shows always bigger number
- -> includes other shared memory objects as well
- Even PG docs recommend to configure on Linux even more than this value
- ... (continuation on the next slide)

# How To Properly Use Huge Pages in PostgreSQL

---

- Set PostgreSQL to use only Huge Pages -> `huge_pages = on`
- Start PostgreSQL -> if it fails, start it again in DEBUG mode
- If it starts, check `shared_memory_size_in_huge_pages` - how many Huge Pages are used
- Check usage of Huge Pages on Linux -> `cat /proc/meminfo |grep -i hugepages`
  - HugePages\_Total: 1024
  - HugePages\_Free: 800
  - HugePages\_Rsvd: 50 (promised pages)
  - -> number of really used Huge Pages: `HugePages_Total - HugePages_Free = 224`
- Shrink `vm.nr_hugepages` to the number of Huge Pages used by PostgreSQL
  - -> Number seems to be stable for given PostgreSQL and Linux versions
  - -> Performance gain seems to be around 10-15% on longer queries
- [Talk: PostgreSQL and Hugepages](#)
- [Cybertec: Huge Pages and PostgreSQL](#)

- Distros enable dynamic **Transparent Huge Pages (THP)** by default
  - Introduced to "democratize" Huge Pages benefits for all applications
  - Can be swapped out -> kernel breaks them back into 4kB pages
  - `khugepaged` - scans memory for contiguous memory regions
  - `kcompactd` - compacts memory / copies pages (on NUMA one per node)
  - `kswapd` - swaps memory
  - Merging or splitting causes latency spikes and locks on memory pages
  - -> THP are not recommended for PostgreSQL due to performance issues
- Why THP are not good for PostgreSQL?
  - During run of query connection allocates memory for processing
  - In smaps marked as "[anonymous]" and "[heap]"
  - This allocation is more or less `work_mem` size -> few MB or dozens of MB
  - -> attempts to merge these pages cause latency spikes and locks on memory pages

- Oracle:
  - Strongly recommends to use Huge Pages
  - -> Uses massive shared memory "System Global Area" (SGA)
  - Requires disable of THP to avoid memory allocation delays
  - -> Can even break HA features due to delayed heartbeat responses
- MySQL / MariaDB:
  - Supports "Large Pages" for InnoDB buffer pool, but configuration is more complex
  - Requires disable of THP especially if "jemalloc" is used
  - (High performance memory allocator)
- MongoDB 8.0+ -> HP not supported, enable THP (v7.0 or earlier - disable THP)
- Redis -> HP not supported (madvise), disable THP - big latency spikes can occur
- Couchbase, Aerospike, ClickHouse -> HP not supported, disable THP

- PostgreSQL implements "separation of concerns" principle
- It still allocates a tiny `SysV interlock segment` on startup
- Used to "advertise" the presence of a running PostgreSQL instance
  - Size is small - typically a few kilobytes
  - It holds `PGShmemHeader` structure - used for inter-process synchronization
  - Defined in `src/include/storage/pg_shmem.h`
  - Contains cluster identity and other control metadata
  - All PG processes attach to it for synchronization
  - Allows instance discovery and avoiding split-brain conflict
  - Used for detection of multiple postmasters

```
## command ipcs shows SysV IPC objects

postgres=# ipcs -m

----- Shared Memory Segments -----
key        shmid      owner      perms      bytes      nattch      status
0x00d295be 0          postgres   600        56          16
```


# Other Shared Memory Objects

- **WAL buffers** - caches WAL records before writing to disk
- **SLRU buffers** - CLOG, MultiXact, CommitTS etc. control structures
- **Lock table** - tracks locks held by transactions
- **Other objects** - various control structures
- Size depends on settings: `max_connections`, `max_locks_per_transaction`
- The `shared_memory_size` parameter reports the size of the main shared memory area (MB)
- Can be exhausted during some operations -> "out of shared memory" error
- Change in settings for shared memory requires restart

```
postgres=# select name, setting, unit from pg_settings where name like '%shared%' order by name;
```

name	setting	unit
dynamic_shared_memory_type	posix	
min_dynamic_shared_memory	0	MB
shared_buffers	16384	8kB
shared_memory_size	179	MB
shared_memory_size_in_huge_pages	90	
shared_memory_type	mmap	
shared_preload_libraries		

- The biggest & most discussed PostgreSQL memory object
  - In-memory copy of tables and indexes data blocks
  - Allocated on server startup, Shared among all connections
  - Blocks kept / evicted based on frequency of access
  - At the beginning exist only as virtual memory



- EXPLAIN ANALYZE shows how successful was query in using shared buffers
  - Buffers: shared hit=X, read=Y, dirtied=Z, written=W
  - shared = blocks found in shared buffers
  - read = blocks not found in shared buffers
  - dirtied = blocks modified in shared buffers
  - written = blocks written to disk
- pg\_stat\_activity - wait\_event\_type / wait\_event
  - BufferPin - waiting for exclusive access to a data buffer
  - -> BufferPin
  - LWLock - light weight lock is held
  - -> BufferContent - Waiting to access a data page in memory (hot pages)
  - -> BufferMapping - Waiting to associate a data block with a buffer (high eviction rate)

- Size is internally stored as number of data pages
  - In source code in `src/backend/utils/init/globals.c`
  - "int NBuffers = 16384;" (= 128MB)
  - NBuffers limit is theoretically INT\_MAX = 2,147,483,647 -> 16 TB
  - -> but capped in code to INT\_MAX /2 = 1,073,741,823 -> 8 TB
  - -> capped in `src/backend/utils/misc/guc_tables.c` - line 2369
- But this size would be very challenging
  - Just Descriptors would take 64 GB
  - Massive TLB trashing (Huge Pages would be required)
  - Most likely extended checkpoints times
  - Non-linear scanning costs in the buffers eviction algorithm


# Shared Buffers Recommended Size

---

- Recommended 25% of the available memory -> but why? And is it still true?
  - Idea is to keep as much data in Linux page cache as possible
  - To avoid big double caching (PostgreSQL shared buffers + Linux page cache)
  - Tests show benefits of larger shared buffers in some cases
  - Even up to 40% - 50% (70%) of total memory
  - But it strongly depends on the workload
  - Content of shared buffers must be relatively static
    - -> Pages must be used many times to justify larger shared buffers
    - -> Frequent big analytical queries are the good candidates here
  - Or the whole database should fit in shared buffers
  - -> But leaving enough memory for connections and OS is important
- Tuning `shared_buffers` for OLTP and data warehouse workloads

# Shared Buffers Structure

- **BufferTag** structure
  - Identifies which disk block is in which buffer
  - Tablespace OID, database OID, Relation file num, Fork num, Block num
  - Fork = physical file: 0=main, 1=free space map, 2=visibility map, 3=init (unlogged)
- **Buffer Descriptors**
  - metadata for each buffer block - locks, dirty flag, usage count, tags
  - max 64 bytes per descriptor -> NBuffers \* 64 bytes
- **BufferDescPadded**
  - Padding to 64 bytes of CPU cache line size
  - Alignment for highly concurrent workloads
  - Avoids false sharing, i.e. unintentional cache invalidation
- **BufferLookupEnt** - hash table for fast lookup
  - Contains BufferTag and associated buffer ID
- **BufferBlocks** array
  - Actual data blocks of tables and indexes -> NBuffers \* 8kB



(Image from the article  
[The Amazing Buffer Tag in PostgreSQL](#))

# Shared Buffers Structure

```
typedef struct buftag
{
    Oid      spcOid;          /* tablespace oid */
    Oid      dbOid;           /* database oid */
    RelFileNumber relNumber;  /* relation file number */
    ForkNumber forkNum;      /* fork number */
    BlockNumber blockNum;    /* blknum relative to begin of reln */
} BufferTag;

typedef struct
{
    BufferTag   key;          /* Tag of a disk page */
    int         id;           /* Associated buffer ID */
} BufferLookupEnt;

typedef struct BufferDesc
{
    BufferTag   tag;          /* ID of page contained in buffer */
    int         buf_id;        /* buffer's index number (from 0) */
    /* state of the tag, containing flags, refcount and usagecount */
    pg_atomic_uint32 state;
    int         wait_backend_pgprocno; /* backend of pin-count waiter */
    int         freeNext;       /* link in freelist chain */
    PgAioWaitRef io_wref;     /* set iff AIO is in progress */
    LWLock      content_lock;  /* to lock access to buffer contents */
} BufferDesc;

#define BUFFERDESC_PAD_TO_SIZE  (SIZEOF_VOID_P == 8 ? 64 : 1)
typedef union BufferDescPadded
{
    BufferDesc  bufferdesc;
    char        pad[BUFFERDESC_PAD_TO_SIZE];
} BufferDescPadded;
```

- Tag Initialization -> BufferTag constructed for desired block
- Hash Computation -> hash code calculated from tag
- Shared Lock Acquisition -> partition lock acquired
  - -> Shared lock allows unlimited concurrent readers
- Lookup -> hash table is searched
- Pinning -> backend retrieves buffer ID from hash table
  - -> increments refcount (18 bits) in buffer descriptor (=pins the buffer), increments usage\_count (max value 5)
  - -> pin must occur while holding partition lock
  - -> without partition lock, buffer could be evicted
- Lock Release -> partition lock released immediately after pinning
  - -> buffer is now protected by refcount
  - -> cannot be evicted until backend unpins it (decrements refcount by 1, keeps usage\_count)

- We assume all buffers are already filled with data
- Tag Initialization -> BufferTag constructed for desired block
- Hash Computation -> hash code calculated from tag
- Shared Lock Acquisition -> partition lock acquired
- Initial Lookup -> hash table is searched - buffer NOT found
- Lock Release -> initial shared lock released
- Victim Selection -> find in Descriptors page with refcount == 0 and usage\_count == 0
- Exclusive Lock Acquisition -> exclusive lock for the new page's partition
- Re-Check -> check if page is still not in the buffer
  - -> other backend might already successfully loaded the page
  - -> if so, exclusive lock released, victim returned to Free list
  - -> if not New BufferTag mapped to BufferID of victim - added to hash table
- Lock Released -> exclusive lock released
- IO Initialization -> IO is initialized

- Clock Sweeping is used to find pages to evict
- Only page with refcount == 0 can be evicted
- -> but if usage\_count > 0, it was recently repeatedly accessed
- -> clock sweep will decrease it -> usage\_count - 1
- Page with refcount == 0 and usage\_count == 0 -> chosen as victim
- -> Actions immediately started for reuse of this buffer
  
- Free List -> list of pages that are not pinned or are empty and can be used
- -> stored in special structure BufferStrategyControl
- Victim -> page found for overwriting because Free list was empty
- Returned page -> page returned to the Free list due to race condition

- `pg_atomic_uint32 state` in `BufferDesc` contains:
  - `Reference Count - refcount` - The lower bits 0-17
  - -> maximum number of concurrent backends that can pin a single buffer
  - -> 18 bits - the maximum value is  $2^{18} - 1 = 262,143$
  - -> But in reality it is capped to value of `MaxBackends`
  - `Usage Count - usage_count` - The bits immediately following the refcount - bits 18-21
  - -> capped to value of 5 - 0=cold, 1-2=warm 3-5=hot
  - `Buffer Flags` - The higher bits - bits 22-31

```
MaxBackends = MaxConnections + autovacuum_worker_slots +
    max_worker_processes + max_wal_senders + NUM_SPECIAL_WORKER_PROCS;

if (MaxBackends > MAX_BACKENDS)
    ereport(ERROR,
        (errcode(ERRCODE_INVALID_PARAMETER_VALUE),
        errmsg("too many server processes configured"), ...
```

# Buffer Descriptor State Buffer Flags


```
#define BM_LOCKED          (1U << 22) /* buffer header is locked */
#define BM_DIRTY           (1U << 23) /* data needs writing */
#define BM_VALID            (1U << 24) /* data is valid */
#define BM_TAG_VALID        (1U << 25) /* tag is assigned */
#define BM_IO_IN_PROGRESS   (1U << 26) /* read or write in progress */
#define BM_IO_ERROR          (1U << 27) /* previous I/O failed */
#define BM_JUST_DIRTIED     (1U << 28) /* dirtied since write started */
#define BM_PIN_COUNT_WAITER (1U << 29) /* have waiter for sole pin */
#define BM_CHECKPOINT_NEEDED (1U << 30) /* must write for checkpoint */
#define BM_PERMANENT         (1U << 31) /* permanent buffer (not unlogged, or init fork) */
```

- `refcount` is managed "manually" (call of Pin/Unpin)
  - -> a bug in the PostgreSQL kernel or an extension can fail to call UnpinBuffer
  - -> such a buffer can never be evicted from shared buffers
  - -> database can run out of shared buffers and stop
- PostgreSQL uses `ResourceOwner` mechanisms to track pins per-transaction
  - -> ResourceOwner cleanup routine forces the unpin
  - -> If a transaction aborts or commits without unpinning
  - -> balance is restored

- Process might need to pin the same buffer multiple times in query
  - -> rare, but possible - nested loop joins / index scans / cursor ops
  - Would lead to lock contention in shared buffers and slow down the query
  - -> connections have PrivateRefCountArray - 8 entries, fits to CPU cache
  - Count is not protected by the buffer lock
- Shared buffer lock is acquired when the PrivateRefCount for a buffer is 0
  - -> either when connection needs block for a first time
  - -> or when refcount was decremented back to 0
  - All other pins / unpins are done in user space
- Mechanism is not simple - for our purposes it is enough to know that it exists

# Partitioned Buffer Locks

- Buffer lookup table is divided into 128 partitions
- Each partition has its own lightweight lock
  - -> no need to lock the whole table
- Lookup or insertion needs to lock only one partition
- Large buffers and many CPUs can lead to lock contention
  - -> implicit constraint on scaling
  - -> some forks increase the number of partitions
  - -> PostgresPro sets NUM\_BUFFER\_PARTITIONS to 512
  - -> OrioleDB tries to combine in-memory and on-disk buffers



- Must be a power of 2 - for efficient bitwise arithmetics
- Value 128 is not arbitrary, but calibrated for trade-off
- Probability of collision for two random tags is  $1/128 = 0.78\%$
- But high number would cause system stalls for some operations
- DROP TABLE, TRUNCATE TABLE, CHECKPOINT might need to lock all part
- -> overhead of acquiring too many locks sequentially
  
- Some experiments suggest improvement with higher values
- -> On machines with hundreds of CPUs and 256+ GB of RAM
- Often combined with bigger data blocks
- -> but this is uncharted territory

- Clock Sweep algorithm expects Zipfian distribution
- -> relatively small working set is accessed frequently
- Sequential scans are a pathological case
- -> huge scans could replace all pages in the buffer
- Solution is the Buffer Access Strategy (BAS) for bulk operations

- **BAS\_BULKREAD** uses a ring buffer of 32 pages (256KB)

- -> can be bigger:

```
ring_size_kb += (BLCKSZ / 1024) * io_combine_limit * effective_io_concurrency;
```

- **BAS\_BULKWRITE** uses a ring buffer of 2048 pages (16MB)

- -> to allow more dirty pages to accumulate before flushing

- **BAS\_VACUUM** uses a ring buffer of 256 pages (2MB)

- -> vacuum is a sequential scan

- Any ring cannot exceed **NBuffers / 8**

# Shared Buffers In PostgreSQL Connection

- Let's look into `/proc/PID/smaps` to see how shared buffers are used in session
  - PC 32GB memory, PostgreSQL 17, shared\_buffers=8GB, effective\_cache\_size=24GB, work\_mem=64MB
  - Connected with psql, connection is newly created, no command issued yet
  - Here is detailed view - showed 42 different `/usr/lib/x86_64-linux-gnu/` libraries
  - And many small regions without paths -> summarized together as `[anonymous]` and `[heap]`
  - Find more in my talk [PostgreSQL Connection Memory Usage](#)

## output of top command														
PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND			
190747	postgres	20	0	8701512	20248	16876	S	0.0	0.1	0:00.00	postgres: postgres postgres	172.18.0.1(40278)	idle	
## python script output - smaps														
Path				Size	Rss	Pss	Pss_Dirty	Shr_Clean	Shr_Dirty	Prv_Clean	Prv_Dirty	Swap	SwapPss	Cnt
<code>/usr/lib/postgresql/16/bin/postgres</code>				9296	4140	1156	75	3792	168	128	52	0	0	5
<code>[anonymous]</code>				1708	660	554	554	0	120	0	540	0	0	21
<code>[heap]</code>				1440	1132	821	821	0	368	0	764	0	0	2
<code>/dev/shm/PostgreSQL.1436672634</code>				1024	132	130	130	0	4	0	128	0	0	1
<code>/dev/shm/PostgreSQL.3104938386</code>				112	4	1	1	0	4	0	0	0	0	1
<code>/dev/zero (deleted)</code>				8624208	10352	5070	5070	0	9780	0	572	0	0	1
<code>/usr/lib/postgresql/16/lib/auto_explain.so</code>				20	8	0	0	0	8	0	0	0	0	5
<code>/usr/lib/postgresql/16/lib/pg_stat_statements.so</code>				44	8	0	0	0	8	0	0	0	0	5
<code>/usr/lib/locale/locale-archive</code>				2980	60	19	0	60	0	0	0	0	0	1
<code>/usr/lib/x86_64-linux-gnu/libffi.so.8.1.2</code>				48	8	0	0	0	8	0	0	0	0	5
<code>/usr/lib/x86_64-linux-gnu/libgpg-error.so.0.33.1</code>				160	8	0	0	0	8	0	0	0	0	5
<code>...</code>														
<code>/SYSV00ce5741 (deleted)</code>				4	0	0	0	0	0	0	0	0	0	1
<code>[stack]</code>				132	36	27	27	0	12	0	24	0	0	1
<code>[vvar]</code>				16	0	0	0	0	0	0	0	0	0	1
<code>[vdso]</code>				8	4	0	0	4	0	0	0	0	0	1
<code>Total</code>				8701512	20380	8553	6841	6356	11760	164	2100	0	0	251

# Why "/dev/zero (deleted)"?

---

- Why shared buffers are mapped as `/dev/zero (deleted)` in `smaps` output?
  - PostgreSQL requests "shared anonymous memory" from OS
  - Using `mmap()` system call with `MAP_ANONYMOUS | MAP_SHARED` flags
  - I.e. "shared memory with anonymous mapping (not backed by any file)" is requested
  - Described in PG code as "anonymous mmap()ed shared memory segment"
  - Hence PG setting `shared_memory_type = 'mmap'`
- How Linux implements this internally?
  - Linux kernel internally instantiates a synthetic file object within `tmpfs`
  - Kernel source code names it "dev/zero" - legacy from older implementations
  - This internal backing file server only as a handle for memory management
  - It is not linked to Virtual File System (VFS) directory tree
  - Therefore it shows up as "(deleted)" in `smaps` output
- Hence shared buffers are mapped as `/dev/zero (deleted)`

# Let's Run Some Heavy Query



- Let's run some heavy aggregations over the table not fitting into memory
- Memory 32 GB, table 38 GB, shared\_buffers 8 GB, work\_mem 64 MB, max\_parallel\_workers\_per\_gather = 0

## top command output after query run																
PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND					
190747	postgres	20	0	8701848	8.2g	8.1g	S	0.0	26.3	0:48.67	postgres: postgres	172.18.0.1(40278)	idle			
## smaps numbers after query run																
Path				Size	Rss			Pss	Pss_Dirty	Shr_Clean	Shr_Dirty	Prv_Clean	Prv_Dirty	Swap	SwapPss	Cnt
/usr/lib/postgresql/16/bin/postgres				9296	6508	3255	79	4176	164	2112	56	0	0	0	5	
[anonymous]				1708	704	598	598	0	120	0	584	0	0	0	21	
[heap]				1776	1516	1208	1208	0	364	0	1152	0	0	0	2	
/dev/shm/				1136	148	143	143	0	8	0	140	980	0	0	2	
/dev/zero (deleted)				8624208	8532008	8443508	8443508	0	143376	0	8388632	0	0	0	1	
/usr/lib/postgresql/16/lib/				64	44	22	8	28	8	0	8	0	0	0	10	
/usr/lib/locale/locale-archive				2980	68	19	0	68	0	0	0	0	0	0	1	
/usr/lib/x86_64-linux-gnu/				60520	5020	1376	167	3556	1284	156	24	0	0	0	205	
/SYSV00ce5741 (deleted)				4	0	0	0	0	0	0	0	0	0	0	1	
[stack]				132	44	44	44	0	0	0	44	0	0	0	1	
[vvar]				16	0	0	0	0	0	0	0	0	0	0	1	
[vdso]				8	4	0	0	4	0	0	0	0	0	0	1	
Total				8701848	8546064	8450173	8445755	7832	145324	2268	8390640	980	0	251		

# But How Much Memory are Connections Really Using?



- I did some playing with multiple sessions with/without parallelism

```
## top command
PID USER      PR  NI    VIRT    RES    SHR S %CPU %MEM     TIME+ COMMAND
190747 postgres  20   0 8701868  8.1g  8.1g S  0.0 26.2  0:48.69 postgres: postgres postgres 172.18.0.1(40278) idle
206119 postgres  20   0 8702808  4.7g  4.7g S  0.0 15.2  0:21.81 postgres: postgres postgres 172.18.0.1(44010) idle
219065 postgres  20   0 8802384  8.2g  8.1g S  0.0 26.6  2:56.64 postgres: postgres postgres 172.18.0.1(52912) idle
228832 postgres  20   0 8709912  3.4g  3.4g S  0.0 11.1  0:11.10 postgres: postgres postgres 172.18.0.1(60090) idle
230390 postgres  20   0 8701340  8.1g  8.1g S  0.0 26.3  0:51.89 postgres: postgres postgres 172.18.0.1(44802) idle

## smaps summaries with /dev/zero
Path          Size      Rss      Pss  Pss_Dirty  Shr_Clean  Shr_Dirty  Prv_Clean  Prv_Dirty      Swap      SwapPss  Cnt
-----
Total for /proc/190747/smaps  8701868  8529172  2196188  2195833  2960  8525332      0      880  61784  1116  256
Total for /proc/206119/smaps  8702808  4928096  1119095  1118712  3120  4924968      0       8  63996  2112  258
Total for /proc/219065/smaps  8802384  8635640  2296848  2295726  5472  8531892      12  98264  64896  4078  252
Total for /proc/228832/smaps  8709912  3609444  798269  795595  8660  3590600      60  10124  60936  100  252
Total for /proc/230390/smaps  8701340  8542712  2202431  2199069  8660  8531564      748  1740  60768  99  251
-----
43618312  34285064  8611831  8613495  34872  34193356      820  19916  312480  17405  1279

## smaps summaries without /dev/zero
Path          Size      Rss      Pss  Pss_Dirty  Shr_Clean  Shr_Dirty  Prv_Clean  Prv_Dirty      Swap      SwapPss  Cnt
-----
Total for /proc/190747/smaps  77660    4416    1291    936    2960    576      0      880  3508  1116  255
Total for /proc/206119/smaps  78600    3708    448     65    3120    580      0       8  5720  2112  257
Total for /proc/219065/smaps  178176   104316   99427   98305   5472    584      12  98248  6620  4078  251
Total for /proc/228832/smaps  85704    19420   12853   10179   8660    576      60  10124  2660  100  251
Total for /proc/230390/smaps  77132    11716    5143    1781   8660    584      748  1724  2492  99  250
-----
499272  169576  118162  110266  34872  2900      820  19984  18300  17405  1274
```


# Why Transparent Huge Pages Are Not Good for PostgreSQL

- Smaps data - single process, table 38 GB, JSONB TOASTed data, work\_mem 8 / 32 / 64 / 128 MB


work\_mem = 8 MB




work\_mem = 32 MB



work\_mem = 64 MB



work\_mem = 128 MB



Max stacked RSS = 17.5 MB  
(2.2x work\_mem)

Sort Method: external merge  
Disk: 307960kB

Max stacked RSS = 61 MB  
(1.9x work\_mem)

Sort Method: external merge  
Disk: 307864kB

Max stacked RSS = 110.6 MB  
(1.7x work\_mem)

Sort Method: external merge  
Disk: 307822kB

Max stacked RSS = 216 MB  
(1.7x work\_mem)

Sort Method: external merge  
Disk: 307792kB

Thank you for your attention!

---



All my slides



Recorded talks

