

Open in app ↗

Medium

 Search Write

# How DuckDB handles data not fitting into memory?



Josef Machytka

4 min read · 2 days ago



In my previous article about DuckDB I described how to use this database as an intelligent ETL tool for PostgreSQL or MySQL. During my testing, I worked with a relatively large (4.5 GB) CSV file from Kaggle, which fit comfortably into the memory of my testing machine. DuckDB's documentation directly mentions that its primary use case is for datasets that fit into memory, noting that while it can handle data spilling to disk, they don't recommend it as a typical use case.

But I was curious about how DuckDB handles data that doesn't fit into memory, so I decided to put it to the test. To simulate this scenario, I used ChatGPT to give me Python code which will generate a large CSV file. I frequently use AI to automate tedious, time-consuming tasks, and this approach saved me considerable time in preparing a realistic testing examples.

## Task 1: Counting Rows

```
D select count(*) from './data/data_100.csv';
44%
```

2/10

```
D select count(*) from './data/data 100.csv';
100%
```

count_star() int64
399682000


```
Run Time (s): real 41.447 user 106.476844 sys 35.288722
```

## Task 2: Summarizing Basic Statistics

For the second task, I asked DuckDB to summarize basic statistics from the file using the SUMMARIZE command.

```
D summarize select * from './data/data_100.csv';
```

DuckDB again utilized multiple cores, this time even more intensively. Eight cores were running at over 98%.



This task took around 79 seconds, clearly documenting DuckDB's efficient parallel processing capabilities.

column name varchar	column type varchar	min varchar	max varchar	approx unique int64	=	q50 varchar	q75 varchar	count int64	null percentage decimal(9,2)
user_id	BIGINT	1	399682000	418759233	=	199903003	299807768	399682000	0.00
first_name	VARCHAR	AAA	zzzzyDjd	322345046	=			399682000	0.00
last_name	VARCHAR	AAA	zzzzuRCU	267610286	=			399682000	0.00
email	VARCHAR	aaa.aao222645931@e_	zzzzzzmg.rplh31199_	389505211	=			399682000	0.00
signup_date	TIMESTAMP	2022-02-11 06:58:44	2024-11-07 09:00:49	8201541	=	2023-06-20 09:41:1_	2024-03-02 15:31:3_	399682000	0.00
last_login	TIMESTAMP	2022-02-11 06:58:49	2024-11-07 09:00:49	8201541	=	2024-05-04 18:53:1_	2024-08-31 19:29:5_	399682000	0.00
is_active	BOOLEAN	false	true	2	=			399682000	0.00
account_balance	DOUBLE	0.0	10000.0	1012389	=	4997.571189812575	7499.3417345063135	399682000	0.00
country_code	VARCHAR	AU	US	11	=			399682000	0.00
favorite_number	BIGINT	1	100	96	=	50	75	399682000	0.00
profile_text	VARCHAR	41rCAFEMWu6bB_	zzzzyCjVWNEr8Qtvc_	424701465	=			399682000	0.00
checksum	VARCHAR	00000005-dac2-4875_	fffffffe-4a11-4d3d_	392541443	=			399682000	0.00
12 rows									12 columns (9 shown)

Run Time (s): real 78.423 user 1279.974541 sys 57.487294

### Task 3: Importing Data into an In-Memory Table

For the third test, I asked DuckDB to import the 100 GB CSV file into an in-memory table. According to the documentation, DuckDB can work with data that doesn't fit into memory, so I wanted to see how it would handle this massive import.

```
D create table data_100gb imported as select * from './data/data_100.csv';
54%
```

This time, I observed DuckDB's approach to handling such a large dataset. Once again, all cores were in use, though less intensively than before. Memory usage grew significantly, and once it reached around 80% of available RAM, the machine began to use swap space.



DuckDB also started creating temporary storage files in the .tmp directory, which at the end reached 63 GB total size.

```
total 65999904
drwxr-xr-x 2 josef josef      4096 Nov 13 06:55 .
drwxrwxr-x 9 josef josef      4096 Nov 13 06:46 ..
-rw-rw-r-- 1 josef josef 1048576000 Nov 13 06:53 duckdb_temp_storage-0.tmp
-rw-rw-r-- 1 josef josef 2138570752 Nov 13 06:53 duckdb_temp_storage-1.tmp
-rw-rw-r-- 1 josef josef 4196925440 Nov 13 06:53 duckdb_temp_storage-2.tmp
-rw-rw-r-- 1 josef josef 8388870144 Nov 13 06:53 duckdb_temp_storage-3.tmp
-rw-rw-r-- 1 josef josef 16779837440 Nov 13 06:54 duckdb_temp_storage-4.tmp
-rw-rw-r-- 1 josef josef 33578024960 Nov 13 06:55 duckdb_temp_storage-5.tmp
-rw-rw-r-- 1 josef josef 1453064192 Nov 13 06:55 duckdb_temp_storage-6.tmp
```

The import operation took almost 190 seconds (3 minutes and 10 seconds). While CPU usage was this time much lower, DuckDB still used masive





After this experiment, my appreciation for DuckDB has grown. It not only handled the 100 GB CSV file but did it efficiently, using parallelism to minimize processing time. DuckDB proves to be a powerful option even for datasets that don't fit into memory. After this test I love DuckDB even more.



Written by Josef Machytka

Edit profile

4 Followers

I work as Professional Service Consultant - PostgreSQL specialist in NetApp Deutschland GmbH, Open Source Services division.

More from Josef Machytka

10.csv';

ix	approx_unique	avg	std	q25	
char	int64	varchar	varchar	varchar	va
	368	200512.6231617008	928.6562707338801	199729	20
	2	1.3918669432239588	0.488167229676373	1	1
	260	193.20352292962244	121.02406459689657	106	12
	146	313.31866609794747	179.77396070679418	104	30
10000	17334				
1000000	924349	47070.754848697165	29145919.848440725	0	0
19678	1767002	127166.90737165902	4571867.453589752	98	71
1402	1394313	32400.302375721876	376908.2117641157	720	24
136306					



Josef Machytka

Josef Machytka

## Using DuckDB as an Intelligent ETL tool for PostgreSQL

There is a lot of hype around DuckDB these days. At one PostgreSQL conference, I even...


Nov 2



## AI Hallucinations are caused by Quantum Pigeons Nesting in...

This is not a new discovery in quantum physics—it is a playful deliberate...

Nov 5



Josef Machytka

### Josef Machytka: Speaker Portfolio


Expert Talks on PostgreSQL, Databases, Data Ingestion and Data Analysis

Nov 3

[See all from Josef Machytka](#)

## Recommended from Medium



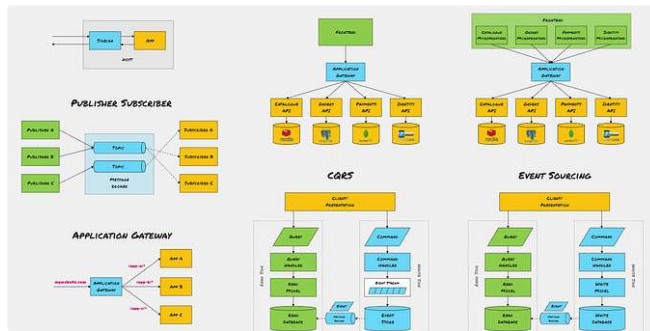


 Harendra

## How I Am Using a Lifetime 100% Free Server

Get a server with 24 GB RAM + 4 CPU + 200 GB Storage + Always Free

★ Oct 26 🖱 4.5K 💬 55 📌 ⋮



 Matt Bentley in Level Up Coding

## My Favourite Software Architecture Patterns

Exploring my most loved Software Architecture patterns and their practical...

★ 2d ago 🖱 1.5K 💬 26 📌 ⋮

### Lists




#### Staff Picks

765 stories · 1437 saves



#### Self-Improvement 101

20 stories · 3013 saves



#### Stories to Help You Level-Up at Work

19 stories · 864 saves

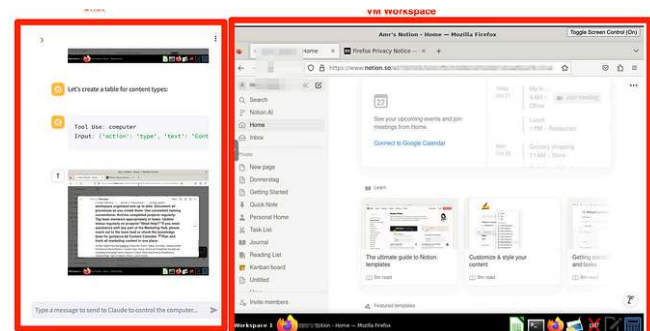


#### Productivity 101

20 stories · 2558 saves



 Ida Silfverskiöld in Towards Data Science



 AI Rabbit in CodeX

## Economics of Hosting Open Source LLMs

Leveraging various deployment options



2d ago



622



8




Oct 27




1.91K



31



I spent 3 hours learning how  **UBER** manages data quality



Vu Trinh in Data Engineer Things

## I spent 3 hours learning how Uber manages data quality.

From the standard to the data quality platform

5d ago



136



1



Nov 4



116



## Has Anthropic Claude just wiped out an entire industry?

If you have been following the news, you may have read about a new feature (or should I ca...



Marek Sirkovský

## New Lock object and history

Discover what's new for locking in .NET 9— even the ancient lock keyword can get an...

See more recommendations