Easy and Intelligent Pivot Tables with DuckDB

Josef Machytka

3 min read · Just now

After exploring the various capabilities of DuckDB in my earlier articles, I want to focus more on its powerful data analytical functionality. In this article, I will delve into the PIVOT command, an amazing powerful tool for creating pivot tables, even from complex datasets.

Pivot tables are a staple for data analysis, but they can be tedious to set up in some systems. DuckDB makes this process both simple and efficient. Let me demonstrate it with a classic “sales” example, inspired by a testing scenario I explored using ChatGPT.

The Example Data

I will use a `sales` table that tracks sales figures for various products, salespeople, and years. Here’s the table structure and some sample data:

```
CREATE TABLE sales (
    id SERIAL PRIMARY KEY,
    salesperson VARCHAR(50),
    product VARCHAR(50),
    year INT,
    sales_amount NUMERIC
);

INSERT INTO sales (salesperson, product, year, sales_amount) VALUES
('Alice', 'Laptop', 2022, 1200),
('Alice', 'Phone', 2022, 800),
('Alice', 'Tablet', 2022, 300),
('Alice', 'Laptop', 2023, 1400),
('Alice', 'Phone', 2023, 900),
('Alice', 'Tablet', 2023, 400),
('Bob', 'Laptop', 2022, 1000),
('Bob', 'Phone', 2022, 600),
('Bob', 'Laptop', 2023, 1100),
('Charlie', 'Tablet', 2022, 500),
('Charlie', 'Laptop', 2022, 1100),
('Charlie', 'Phone', 2022, 700),
('Charlie', 'Tablet', 2023, 600),
('Charlie', 'Laptop', 2023, 1200),
('Charlie', 'Phone', 2023, 800);
```

Pivoting in PostgreSQL: The Hard Way

Creating a pivot table in standard PostgreSQL is of course perfectly possible but involves a fair amount of manual work. For example, let's pivot the `sales` table to display sales amounts for each salesperson by product and year. Here's the query:

```
SELECT
    salesperson,
    SUM(sales_amount) FILTER (WHERE product = 'Laptop' AND year = 2022) AS Lapto
    SUM(sales_amount) FILTER (WHERE product = 'Phone' AND year = 2022) AS Phone_
    SUM(sales_amount) FILTER (WHERE product = 'Tablet' AND year = 2022) AS Table
    SUM(sales_amount) FILTER (WHERE product = 'Laptop' AND year = 2023) AS Lapto
    SUM(sales_amount) FILTER (WHERE product = 'Phone' AND year = 2023) AS Phone_
    SUM(sales_amount) FILTER (WHERE product = 'Tablet' AND year = 2023) AS Table
```

```
FROM
    sales
GROUP BY
    salesperson
ORDER BY
    salesperson;
```

This works but requires explicitly specifying every product and year combination in the query. If people later add new products or years, we must rewrite the query to account for those changes. Clearly, this approach doesn't scale well and involves too much manual effort.

Pivoting in DuckDB: Let's Make Pivot Fun Again

DuckDB offers a simpler and more flexible solution with its internal **PIVOT** command. By attaching a remote PostgreSQL database, we can process the same data in a far more elegant way. Adjusting pivot criteria is effortless, making it an incredibly user-friendly tool for dynamic data exploration.

```
D pivot pg.sales on product using sum(sales_amount) group by salesperson order by salesperson;
```

salesperson varchar	Laptop double	Phone double	Tablet double
Alice	2600.0	1700.0	700.0
Bob	2100.0	600.0	
Charlie	2300.0	1500.0	1100.0

```
D pivot pg.sales on (product,year) using sum(sales_amount) group by salesperson order by salesperson;
```

salesperson varchar	(Laptop, 2022) double	(Laptop, 2023) double	(Phone, 2022) double	(Phone, 2023) double	(Tablet, 2022) double	(Tablet, 2023) double
Alice	1200.0	1400.0	800.0	900.0	300.0	400.0
Bob	1000.0	1100.0	600.0			
Charlie	1100.0	1200.0	700.0	800.0	500.0	600.0

```
D pivot pg.sales on (year,product) using sum(sales_amount) group by salesperson order by salesperson;
```

salesperson varchar	(2022, Laptop) double	(2022, Phone) double	(2022, Tablet) double	(2023, Laptop) double	(2023, Phone) double	(2023, Tablet) double
Alice	1200.0	800.0	300.0	1400.0	900.0	400.0
Bob	1000.0	600.0		1100.0		
Charlie	1100.0	700.0	500.0	1200.0	800.0	600.0

```
D pivot pg.sales on (year) using sum(sales_amount) group by salesperson order by salesperson;
```

salesperson varchar	2022 double	2023 double
Alice	2300.0	2700.0
Bob	1600.0	1100.0
Charlie	2300.0	2600.0

Conclusion

DuckDB transforms the process of creating pivot tables into a seamless experience, saving us both time and effort. As shown in these examples, its built-in capabilities outperform traditional SQL methods for pivoting. Whether we are dealing with a simple dataset or a complex one, DuckDB makes data analysis not just easier but also more enjoyable.

Duckdb | Pivot Tables | Pivoting | Postgresql

Written by Josef Machytka

16 Followers · 3 Following

Edit profile

I work as Professional Service Consultant - PostgreSQL specialist in NetApp Deutschland GmbH, Open Source Services division.

More from Josef Machytka

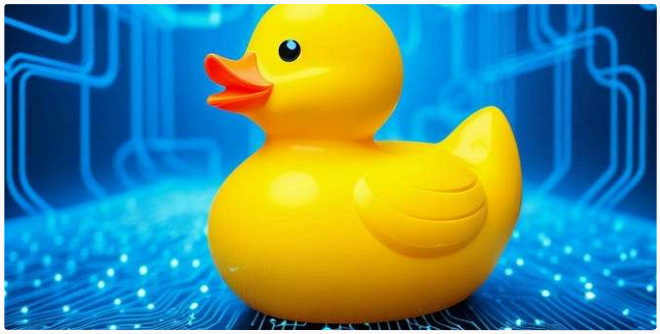

Josef Machytka

How DuckDB handles data not fitting into memory?

In my previous article about DuckDB I described how to use this database as an...

Nov 13 7

Josef Machytka

Extending DuckDB ETL Capabilities with Python

DuckDB has recently become my go-to solution for small ETL tasks. It is an...

Nov 24 12

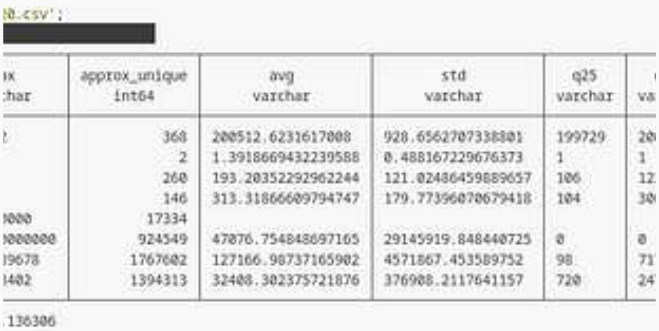

ix :har	approx_unique int64	avg varchar	std varchar	q25 varchar	va
	368	200512.6231617008	928.6562707338801	199729	20
	2	1.3918669432239588	0.488167229676373	1	1
	260	193.20352292962244	121.02486459889657	106	12
	146	313.31866609794747	179.77396070679418	104	30
000	17334				
000000	924549	47876.754848697165	29145919.848440725	0	0
9678	1767602	127166.98737165902	4571867.453589752	98	71
402	1394313	32408.302375721876	376908.2117641157	720	24

136306

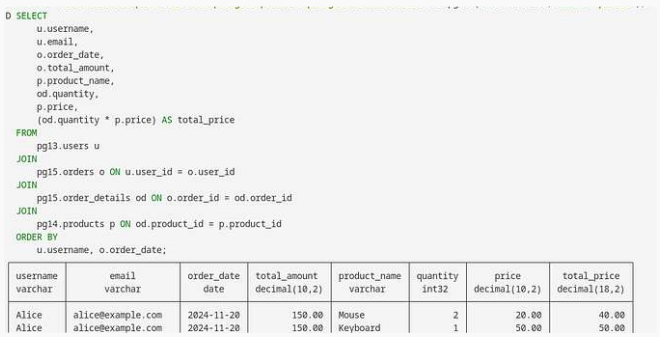

```
D SELECT
    u.username,
    u.email,
    o.order_date,
    o.total_amount,
    p.product_name,
    od.quantity,
    p.price,
    (od.quantity * p.price) AS total_price
FROM
    pg13.users u
JOIN
    pg15.orders o ON u.user_id = o.user_id
JOIN
    pg15.order_details od ON o.order_id = od.order_id
JOIN
    pg14.products p ON od.product_id = p.product_id
ORDER BY
    u.username, o.order_date;
```

username varchar	email varchar	order_date date	total_amount decimal(10,2)	product_name varchar	quantity int32	price decimal(10,2)	total_price decimal(18,2)
Alice	alice@example.com	2024-11-20	150.00	Mouse	2	20.00	40.00
Alice	alice@example.com	2024-11-20	150.00	Keyboard	1	50.00	50.00

Josef Machytka

Using DuckDB as an Intelligent ETL tool for PostgreSQL

There is a lot of hype around DuckDB these days. At one PostgreSQL conference, I even...

Nov 2 6

Josef Machytka

Easy Cross-Database Selects with DuckDB

DuckDB was created with simplicity and ease of use in mind. In my previous article, I...

Nov 26 2

See all from Josef Machytka

Recommended from Medium

In T3CH by Jesus L. Monroy

Supercharge Your SQL Analysis with Python and DuckDB

Rethinking Data Analysis: A New Paradigm

5d ago 90

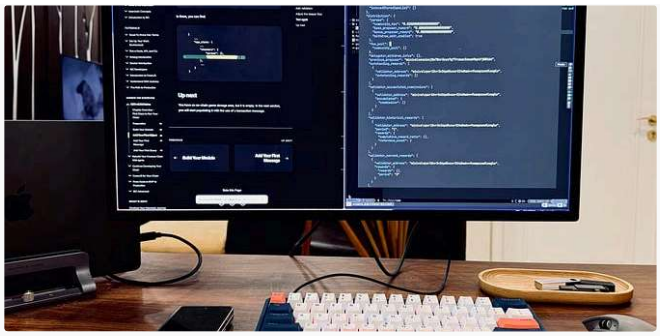

In Code Like A Girl by Chris Ng

I Removed VS Code, and My Productivity Improved...

A Journey from Clunky to Efficient: How The New IDE Transformed My Workflow

Nov 26 553 7

Lists

Staff picks

780 stories · 1488 saves

Stories to Help You Level-Up at Work

19 stories · 888 saves

Self-Improvement 101

20 stories · 3114 saves

Productivity 101

20 stories · 2623 saves

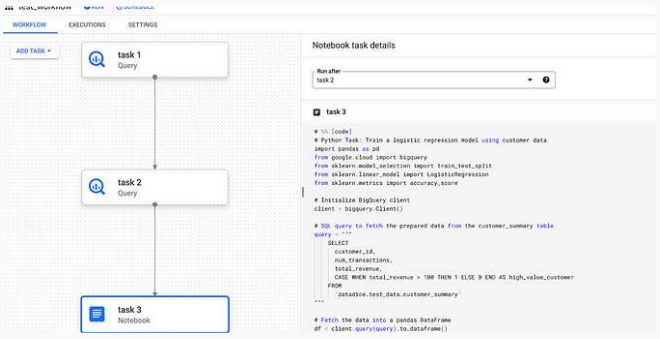

datadice

In Towards Data Science by Eirik Berge, PhD

Introducing BigQuery Workflows

A New Era of Data Orchestration. By Fassil S. Yehuala

Nov 27 108

Use Tablib to Handle Simple Tabular Data in Python

Sometimes a Shallow Abstraction is more Valuable than Performance

Nov 27 211 1

In Level Up Coding by Md Arman Hossen

The Lazy Developer's Guide to Python Automation (and Why...

It was another late night at the office. The soft glow of my monitor illuminated the scattere...

Nov 7 251 1

Thibaut Gourdel

Data Cleaning Tools in 2024

Data cleaning tools: Code-based (pandas, dbt), Open-source (Amphi, OpenRefine), and...

Nov 26 25

See more recommendations