

26 JUNE 2024

Prague PostgreSQL Developers Day 2024

On Wednesday, June 5th, I attended the [Prague PostgreSQL Developers Day 2024](#)[↗]. It is the premier PostgreSQL conference in the Czech Republic, and this year marked its 16th iteration. The event was held in the modern and convenient environment of the Czech Technical University and was attended by almost 270 participants.

During the conference, I presented my talk, "[GIN, BTREE, GIN, GIST, and BTREE Indexes on JSONB Data](#)[↗]." This talk summarized the current findings of my project at NetApp (credativ), which I initiated to deepen our understanding of these indexing methods and their performance results. Our goal is to provide relevant and valuable solutions to our customers, who often struggle with implementing JSONB columns and operations into their applications and find the available information insufficient. Even existing AI models fall short as they rely on the same limited publicly available data. The project is focused on JSONB data; however, the results have already shown applicability beyond just this type of data. The positive reactions from the audience indicated that my presentation was well-received. The conference is a bilingual event, with presentations in both Czech and English. Being from Czechia, I delivered my talk in Czech but used English slides.

The conference also featured seven other insightful talks throughout the day. The [first talk by Jan Karremans](#)[↗] from Cybertec delved into the [CloudNativePG operator](#)[↗], which implements PostgreSQL to run on Kubernetes. The second talk by Jakub Zemanek from initMax provided a detailed guide on [configuring PostgreSQL accounts](#)[↗] based on Active Directory using Ansible and the Idap2pg program. Gülçin Yıldırım Jelínek from EDB followed with a [practical example of implementing vector storage](#)[↗] of AI embeddings in PostgreSQL. I also really enjoyed the engaging and informative talk by Boriss Mejias from EDB, who explained [different types of consistency](#)[↗] and ways to ensure them in PostgreSQL using synchronous or asynchronous replication, and colored his talk with practical examples from a big musical event, where payments from multiple sources must be quickly and safely distributed throughout the system. Pavel Stehule offered a deep analysis of the pros and cons of the existing [cost-based optimizer in PostgreSQL](#)[↗] (slides in Czech lang). Pavlo Golub from Cybertec delivered an informative session on programming different [PostgreSQL operations with the Go language](#)[↗]. In the final talk, Ales Zeleny presented an in-depth analysis of the features of two extensions focused on monitoring query performance in PostgreSQL – [pg_stat_statements and pg_stat_monitor](#)[↗].

The conference was really well-organized, and I extend my gratitude to Tomas Vondra and the other organizers for their hard work. Overall, it was very successful event, filled with valuable insights, engaging discussions, and networking opportunities. I look forward to applying the knowledge gained here and I am eager for future editions of the P2D2 conference.

Useful links:

1. [Prague PostgreSQL Developer Day 2024 talks](#)[↗] – summaries with links to slides (English / Czech)
2. [Prague PostgreSQL Developer Day 2024 conference recap – Tomas Vondra](#)[↗] (English)

Photos by Tomas Vondra (EDB):



ABOUT THE AUTHOR

Josef Machytka

[View posts](#) 