

[Open in app](#)

Medium

 Search Write

Using DuckDB as an Intelligent ETL tool for PostgreSQL



Josef Machytka

4 min read · 1 day ago



There is a lot of hype around DuckDB these days. At one PostgreSQL conference, I even saw a large poster comparing DuckDB with PostgreSQL, presenting DuckDB as a superior tool. Having worked with various databases over many years, I was really curious about this new tool and decided to test it deeper. My experiments confirmed that DuckDB is indeed a valuable tool, filling gaps in areas where other databases fall short or show limited focus.

I love DuckDB. It is already a powerful tool and has the potential to be a game-changer for specific use cases within the next year. However, by my opinion some articles rather exaggerate its current capabilities. It only reached production maturity in June of this year and is currently at version 1.1.2, which has addressed many bugs. Some compelling extensions, which I would love to use and recommend to clients, remain marked as “experimental.” But hopefully in the next six to twelve months, DuckDB will truly shine. That said, it already supports significant use cases today. So let us dive into one.

I have always worked with big data, including importing large external data files into PostgreSQL and MySQL. A few years ago, ad-hoc import of a CSV file several gigabytes in size into these databases was almost a celebrated achievement. Even now, it requires manual checks and preparations, especially in PostgreSQL. Existing external tools are not always user friendly and most of them are really slow. This is why I am thrilled that DuckDB can greatly streamline this process..

What sets DuckDB apart when working with external data files is its simplicity. There is no need to create database objects manually or even check the file's structure beforehand. DuckDB implements external data sources via table functions API, enabling us to select data from external files directly within the ``FROM`` clause of a query.

But what is the most amazing part about DuckDB when it comes to external data files, is the absolute simplicity of usage. No need to create any object manually, no need to even check structure of the file. DuckDB implements all external data sources through table functions API. Which means we can directly select data from the external file in the FROM clause of the query.

For my tests, I used a 100-million-row dataset from Kaggle. This 4.5 GB CSV file contains 113 million rows of Japanese trade statistics dating back to 1988. The structure is loosely described as eight INT columns, and the file lacks a header. In the past, importing this type of data would typically involve several minutes of manual setup and problem-solving.

DuckDB, however, handled this file effortlessly. The data selection was as simple as:

```
SELECT * FROM '/data/custom_1988_2020.csv';
```

To inspect the file structure, I used:

```
DESCRIBE SELECT * FROM '/data/custom_1988_2020.csv';
```

To my surprise, DuckDB immediately identified an issue with the dataset: numeric values in column 5 were formatted as strings, causing conversion errors when attempting to import the data as eight INT columns, as described on Kaggle. DuckDB flagged this issue in just 0.03 seconds.

```
D DESCRIBE SELECT * FROM '/data/custom_1988_2020.csv';
```

column_name varchar	column_type varchar	null varchar	key varchar	default varchar	extra varchar
column0	BIGINT	YES			
column1	BIGINT	YES			
column2	BIGINT	YES			
column3	BIGINT	YES			
column4	VARCHAR	YES			
column5	BIGINT	YES			
column6	BIGINT	YES			
column7	BIGINT	YES			

```
Run Time (s): real 0.032 user 0.018639 sys 0.013101
```

Another amazing feature is DuckDB's ability to create basic column statistics with minimal effort. Using:

```
SUMMARIZE SELECT * FROM '/data/custom_1988_2020.csv';
```

I received comprehensive statistics within approximately 12 seconds. Doing this in other databases would require writing custom analytical queries first, which makes DuckDB a clear winner for quick data profiling.

```
D summarize select * from '/data/custom_1988_2020.csv';
100%
```

column_name varchar	column_type varchar	min varchar	max varchar	approx_unique int64	avg varchar	std varchar	q25 varchar	q50 varchar	q75 varchar	count int64	null_percentage decimal(9,2)
column0	BIGINT	198801	202012	360	200512.6231617000	928.6562707338801	199729	200544	201314	113607322	0.00
column1	BIGINT	1	2	2	1.3918669432239588	0.488167229676373	1	1	2	113607322	0.00
column2	BIGINT	103	703	260	193.20352292962244	121.02486459889657	106	123	222	113607322	0.00
column3	BIGINT	100	908	146	313.31866609794747	179.77396070679418	104	300	428	113607322	0.00
column4	VARCHAR	000000011	970000000	17334						113607322	0.00
column5	BIGINT	0	12550000000	924549	47076.754048697165	29145919.848440725	0	0	8	113607322	0.00
column6	BIGINT	0	1085789670	1767602	127166.98737165902	4571067.453589752	98	717	6760	113607322	0.00
column7	BIGINT	30	183278402	1394313	32408.302375721876	376988.2117641157	720	2478	10072	113607322	0.00

```
Run Time (s): real 11.671 user 219.519242 sys 1.136306
```

Transferring data from DuckDB to PostgreSQL is straightforward and involves just two steps. First, you attach the remote PostgreSQL database and specify a remote schema name. Then, you import the data using “CREATE TABLE in PostgreSQL AS SELECT * FROM data file”:

```
D ATTACH 'dbname=duckdb_test user=postgres host=postgres_container port=5432 password=postgres' AS pg_duckdb_test (TYPE postgres, SCHEMA 'public');
Run Time (s): real 0.627 user 0.394838 sys 0.082223
D CREATE TABLE pg_duckdb_test.custom_1988_2020 AS SELECT * FROM '/data/custom_1988_2020.csv';
100%
Run Time (s): real 88.547 user 36.957726 sys 2.746653
```

Credentials can be provided directly within the `ATTACH` command or managed through DuckDB's `CREATE SECRET` function. And in PostgreSQL we can see our new table:

```
duckdb_test=# \dS+ custom_1988_2020
```

Table "public.custom_1988_2020"									
Column	Type	Collation	Nullable	Default	Storage	Compression	Stats target	Description	
column0	bigint				plain				
column1	bigint				plain				
column2	bigint				plain				
column3	bigint				plain				
column4	character varying				extended				
column5	bigint				plain				
column6	bigint				plain				
column7	bigint				plain				

Access method: heap

```
duckdb_test=# select count(*) from custom_1988_2020;
 count
-----
113607322
(1 row)
```

And that's it! That's the entire process for importing a 4.5 GB CSV file into PostgreSQL. The same method works for importing data into MySQL. This also works seamlessly with JSON and Parquet files, and hopefully, extensions for other data formats will be available soon.

After my tests, I really love DuckDB and I am convinced that it is an essential addition to the toolkit of any data scientist, data analyst, or database administrator. However, it is not a competitor to “big” databases; instead, it complements them by addressing weaknesses in areas that are often overlooked or underestimated.

[Postgresql](#)[Duckdb](#)[Csv Import](#)[Data Analysis](#)



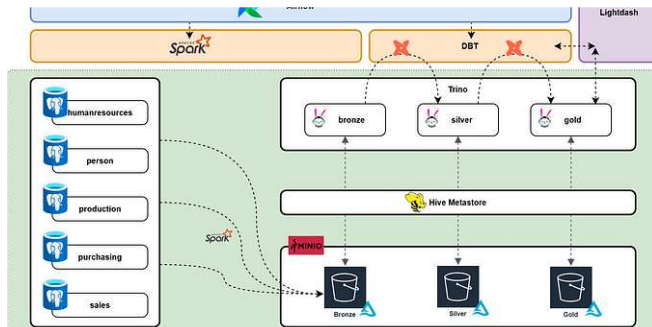
Written by Josef Machytka

[Edit profile](#)

0 Followers

I work as Professional Service Consultant - PostgreSQL specialist in NetApp Deutschland GmbH, Open Source Services division.

Recommended from Medium



Thanh Enc in Data Engineer Things

PoC Data Platform project utilizing modern data stack (Airflow, Spar...

PoC Data Platform is an innovative project designed to demonstrate how data from...

Oct 8 150 1



Nikita Shpilevoy

Advanced SQL Techniques and Complex Queries in PostgreSQL

PostgreSQL stands out as a open source RDBMS known for its adherence, to SQL...

Oct 1 5



Lists



Practical Guides to Machine Learning

10 stories · 1995 saves



ChatGPT prompts

50 stories · 2169 saves




Staff Picks

755 stories · 1417 saves



Natural Language Processing

1789 stories · 1394 saves



Anita Shaw

Optimizing PostgreSQL Queries

Optimizing PostgreSQL Queries: Techniques and Examples

Sep 27 🖱️ 1

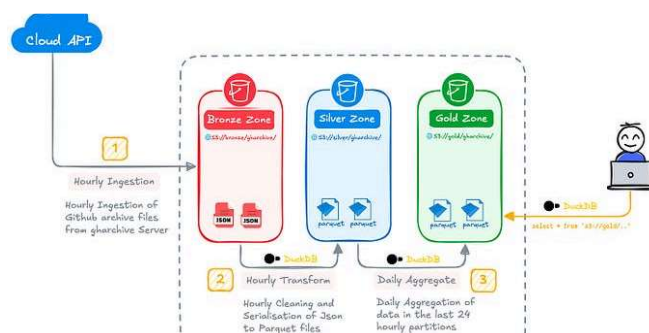


Abdur Rahman in Stackademic

Python is No More The King of Data Science


5 Reasons Why Python is Losing Its Crown

Oct 23 🖱️ 2K 💬 16



Alireza Sadeghi

Building a High-Performance Data Pipeline Using DuckDB



Stefan Krawczyk

Open Source Python Data Lineage with OpenLineage and Hamilton

Using DuckDB to Serialise, Transform, and Aggregate Data in Data Lakes

Oct 20  250  2



Data lineage for SQL is basically a solved problem. For python based workflows, that i...

Sep 25  34  1



See more recommendations